

Student's name:

Class:

**Class Exercises**  
**Module: Distributed Systems**  
**Chapter 6: Synchronization (2/2)**

Question 1: What is a mutual exclusion algorithm in a distributed system?

Question 2: What is the drawback of the centralized algorithm for the mutual exclusion?

Question 3: What is the drawback of the distributed algorithm for the mutual exclusion?

Question 4: Propose a solution for the problem of lost token in Token Ring mutual exclusion algorithm. (hint: you can improve/modify the original algorithm).

Question 5: A system of 8 nodes ( $P_0$  to  $P_7$ ) applies the Bully election algorithm. There are two broken nodes:  $P_4$  and  $P_7$ . The node  $P_3$  starts the election. How many messages does the system need to vote the coordinator?

Question 6: The nodes of a system are assigned an ID from 1 to  $N$ . Each node maintains a status table that consists of states of others nodes with two states: *Running* and *Broken*. An election algorithm is described as follows:

When a node  $P_i$  detects that the current coordinator is broken, it will look in the status table and send *ELECTION* message to the node whose ID just below the broken coordinator's ID. If the node that receives *ELECTION* still runs normally, it sends *OK* to  $P_i$  and then broadcasts the *COORDINATOR* message to all other nodes to inform that it is the new coordinator. If the node that receives *ELECTION* is broken,  $P_i$  continues to repeat the previous step with the nodes that have greater ID than himself (the information retrieved in the status table) until  $P_i$  receives *OK*, or the ID drops to its own ID. If the ID drops to its own ID,  $P_i$  knows that it becomes the new coordinator and broadcasts *COORDINATOR* message to the whole system.

- a) How many messages do we need to vote the coordinator?
- b) Propose an additional part for the above algorithm to handle the problem when a broken node comes back to running state. How many messages do we need for this recovery procedure? (hint: consider 2 cases where this node ID is greater and smaller than the current coordinator ID).

(\* note: the answer related to the number of messages consists only two variables  $N$  and  $i$ . Do not put other variables into the result.)

Question 7: A mutual exclusion algorithm is described as follows:

Consider a system having  $n$  processes:  $P_1, P_2, \dots, P_n$ . There is a shared resource called SR (Shared Resource). Each process  $P_i$  maintains a queue  $queue_i$  to store requests that have not yet been executed.

When the process  $P_i$  wants to access the SR, it will broadcast a *REQUEST* message  $(ts_i, i)$  to all other processes, and store that message in its  $queue_i$  in which  $ts_i$  is timestamp of request.

When a process  $P_j$  receives a *REQUEST* $(ts_i, i)$  from the  $P_i$  process, it takes that request into its queue ( $queue_j$ ) and sends back to  $P_i$  a *REPLY* message.

The  $P_i$  process will allow itself to use SR when it checks that its request is located in the queue  $queue_i$  and other requests have a greater timestamp than its own.

After using SR, the process  $P_i$  deletes its request from the queue and broadcasts the *RELEASE* message for all other processes. When the process  $P_j$  receives a *RELEASE* message from  $P_i$ , it will delete the request of  $P_i$  in its queue.

- a) How many messages does the system need to successfully let a process use SR?
- b) There is an improvement the above algorithm described as follows: after  $P_j$  sends a *REQUEST* to other processes, it receives a *REQUEST* message from  $P_i$ , if it finds the its *REQUEST* timestamp is greater than the *REQUEST* timestamp of  $P_i$ , it won't send a *REPLY* message to  $P_i$  anymore. Does this improvement work? In applying this improvement, how many messages does the system need to successfully let a process use SR? Explain it.