DISTRIBUTED SYSTEMS
PRINCIPLES AND PARADIGMS

SECOND EDITION

PROBLEM SOLUTIONS

ANDREW S. TANENBAUM

MAARTEN VAN STEEN

Vrije Universiteit
Amsterdam, The Netherlands

SOLUTIONS TO CHAPTER 1 PROBLEMS

1. **Q:** An alternative definition for a distributed system is that of a collection of independent computers providing the view of being a *single system*, that is, it is completely hidden from users that there even multiple computers. Give an example where this view would come in very handy.

   A: What immediately comes to mind is parallel computing. If one could design programs that run without any serious modifications on distributed systems that appear to be the same as nondistributed systems, life would be so much easier. Achieving a single-system view is by now considered virtually impossible when performance is in play.

2. **Q:** What is the role of middleware in a distributed system?

   A: To enhance the distribution transparency that is missing in network operating systems. In other words, middleware aims at improving the single-system view that a distributed system should have.

3. **Q:** Many networked systems are organized in terms of a back office and a front office. How does organizations match with the coherent view we demand for a distributed system?

   A: A mistake easily made is to assume that a distributed system as operating in an organization, should be spread across the entire organization. In practice, we see distributed systems being installed along the way that an organization is split up. In this sense, we could have a distributed system supporting back-office procedures and processes, as well as a separate front-office system. Of course, the two may be coupled, but there is no reason for letting this coupling be fully transparent.

4. **Q:** Explain what is meant by (distribution) transparency, and give examples of different types of transparency.

   A: Distribution transparency is the phenomenon by which distribution aspects in a system are hidden from users and applications. Examples include access transparency, location transparency, migration transparency, relocation transparency, replication transparency, concurrency transparency, failure transparency, and persistence transparency.

5. **Q:** Why is it sometimes so hard to hide the occurrence and recovery from failures in a distributed system?

   A: It is generally impossible to detect whether a server is actually down, or that it is simply slow in responding. Consequently, a system may have to report that a service is not available, although, in fact, the server is just slow.

6. **Q:** Why is it not always a good idea to aim at implementing the highest degree of transparency possible?

   A: Aiming at the highest degree of transparency may lead to a considerable loss of performance that users are not willing to accept.

7. **Q:** What is an open distributed system and what benefits does openness provide?

   A: An open distributed system offers services according to clearly defined rules. An open system is capable of easily interoperating with other open systems but also allows applications to be easily ported between different implementations of the same system.

8. **Q:** Describe precisely what is meant by a scalable system.

   A: A system is scalable with respect to either its number of components, geographical size, or number and size of administrative domains, if it can grow in one or more of these dimensions without an unacceptable loss of performance.

9. **Q:** Scalability can be achieved by applying different techniques. What are these techniques?

   A: Scaling can be achieved through distribution, replication, and caching.

10. **Q:** Explain what is meant by a virtual organization and give a hint on how such organizations could be implemented.

    A: A virtual organization (VO) defines a group of users/applications that have access to a specified group of resources, which may be distributed across many different computers, owned by many different organizations. In effect, a VO defines who has access to what. This also suggests that the resources should keep an account of foreign users along with their access rights. This can often be done using standard access control mechanisms (like the rwx bits in UNIX), although foreign users may need to have a special account. The latter complicates matters considerably.

11. **Q:** When a transaction is aborted, we have said that the world is restored to its previous state, as though the transaction had never happened. We lied. Give an example where resetting the world is impossible.

    A: Any situation in which physical I/O has occurred cannot be reset. For example, if the process has printed some output, the ink cannot be removed from the paper. Also, in a system that controls any kind of industrial process, it is usually impossible to undo work that has been done.

12. **Q:** Executing nested transactions requires some form of coordination. Explain what a coordinator should actually do.

    A: A coordinator need simply ensure that if one of the nested transactions aborts, that all other subtransactions abort as well. Likewise, it should

coordinate that all of them commit when each of them can. To this end, a nested transaction should wait to commit until it is told to do so by the coordinator.

13. **Q:** We argued that distribution transparancy may not be in place for pervasice systems. This statement is not true for all types of transparencies. Give an example.

    A: Think of migration transparency. In mnay pervasive systems, components are mobile and will need to re-establish connections when moving from one access point to another. Preferably, such handovers should be completely transparent to the user. Likewise, it can be argued that many other types of transparencies should be supported as well. However, what should not be hidden is a user is possibly accessing resources that are directly coupled to the user's current environment.

14. **Q:** We already gave some examples of distributed pervasive systems: home systems, electronic health-care systems, and sensor networks. Extend this list with more examples.

    A: There are quite a few other examples of pervasive systems. Think of large-scale wireless mesh networks in cities or neighborhoods that provide services such as Internet access, but also form the basis for other services like a news system. There are systems for habitat monitoring (as in wildlife resorts), electronic jails by which offenders are continuously monitored, large-scale integrated sports systems, office systems deploying active badges to know about the whereabouts of their employees, and so on.

15. **Q:** Sketch a design for a home system consisting of a separate media server that will allow for the attachment of a wireless client. The latter is connected to (analog) audio/video equipment and transforms the digital media streams to analog output. The server runs on a separate machine, possibly connected to the Internet, but has no keyboard and/or monitor connected.

SOLUTIONS TO CHAPTER 2 PROBLEMS

1. **Q:** If a client and a server are placed far apart, we may see network latency dominating overall performance. How can we tackle this problem?

   A: It really depends on how the client is organized. It may be possible to divide the client-side code into smaller parts that can run separately. In that case, when one part is waiting for the server to respond, we can schedule another part. Alternatively, we may be able to rearrange the client so that it can do other work after having sent a request to the server. This last solution effectively replaces the synchronous client-server communication with asynchronous one-way communication.

2. Q: What is a three-tiered client-server architecture?

A: A three-tiered client-server architecture consists of three logical layers, where each layer is, in principle, implemented at a separate machine. The highest layer consists of a client user interface, the middle layer contains the actual application, and the lowest layer implements the data that are being used.

3. Q: What is the difference between a vertical distribution and a horizontal distribution?

A: Vertical distribution refers to the distribution of the different layers in a multitiered architectures across multiple machines. In principle, each layer is implemented on a different machine. Horizontal distribution deals with the distribution of a single layer across multiple machines, such as distributing a single database.

4. Q: Consider a chain of processes P_1, P_2, ..., P_n implementing a multitiered client-server architecture. Process P_i is client of process P_{i+1}, and P_i will return a reply to P_{i-1} only after receiving a reply from P_{i+1}. What are the main problems with this organization when taking a look at the request-reply performance at process P_1?

A: Performance can be expected to be bad for large n. The problem is that each communication between two successive layers is, in principle, between two different machines. Consequently, the performance between P_1 and P_2 may also be determined by $n-2$ request-reply interactions between the other layers. Another problem is that if one machine in the chain performs badly or is even temporarily unreachable, then this will immediately degrade the performance at the highest level.

5. Q: In a structured overlay network, messages are routed according to the topology of the overlay. What is an important disadvantage of this approach?

A: The problem is that we are dealing only with *logical* paths. It may very well be the case that two nodes A and B which are neighbors in the overlay network are physically placed far apart. As a consequence, the logically short path between A and B may require routing a message along a very long path in the underlying physical network.

6. Q: Consider the CAN network from Fig. 2-0. How would you route a message from the node with coordinates (0.2,0.3) to the one with coordinates (0.9,0.6)?

A: There are several possibilities, but if we want to follow the shortest path according to a Euclidean distance, we should follow the route (0.2,0.3) → (0.6,0.7) → (0.9,0.6), which has a distance of 0.882. The alternative route (0.2,0.3) → (0.7,0.2) → (0.9,0.6) has a distance of 0.957.

7. Q: Considering that a node in CAN knows the coordinates of its immediate neighbors, a reasonable routing policy would be to forward a message to the closest node toward the destination. How good is this policy?

A: In our example from the previous question, it can already be seen that it need not lead to the best route. If node (0.2,0.3) follows this policy for the message destined for node (0.9,0.6), it would send it off to node (0.7,0.2).

8. Q: Consider an unstructured overlay network in which each node randomly chooses c neighbors. If P and Q are both neighbors of R, what is the probability that they are also neighbors of each other?

A: Consider a network of N nodes. If each node chooses c neighbors at random, then the probability that P will choose Q, or Q chooses P is roughly $2c / (N-1)$.

9. Q: Consider again an unstructured overlay network in which every node randomly chooses c neighbors. To search for a file, a node floods a request to its neighbors and requests those to flood the request once more. How many nodes will be reached?

A: An easy upper bound can be computed as $c \times (c-1)$, but in that case we ignore the fact that neighbors of node P can be each other's neighbor as well. The probability q that a neighbor of P will flood a message only to nonneighbors of P is 1 minus the probability of sending it to at least one neighbor of P:

$$q = 1 - \sum_{k=1}^{c-1} \binom{c-1}{k} \left(\frac{c}{N-1} \right)^k \left(1 - \frac{c}{N-1} \right)^{c-1-k}$$

In that case, this flooding strategy will reach $c \times q \, (c-1)$ nodes. For example, with $c = 20$ and $N = 10,000$, a query will be flooded to 365.817 nodes.

10. Q: Not every node in a peer-to-peer network should become superpeer. What are reasonable requirements that a superpeer should meet?

A: In the first place, the node should be highly available, as many other nodes rely on it. Also, it should have enough capacity to process requests. Most important perhaps is that fact that it can be trusted to do its job well.

11. Q: Consider a BitTorrent system in which each node has an outgoing link with a bandwidth capacity B_{out} and an incoming link with bandwidth capacity B_{in}. Some of these nodes (called seeds) voluntarily offer files to be downloaded by others. What is the maximum download capacity of a BitTorrent client if we assume that it can contact at most one seed at a time?

A: We need to take into account that the outgoing capacity of seeding nodes needs to be shared between clients. Let us assume that there are S seeders and N clients, and that each client randomly picks one of the seeders. The joint outgoing capacity of the seeders is $S \times B_{out}$, giving each of the clients $S \times B_{out} / N$ immediate download capacity. In addition, if clients help each

other, each one of them will be able to download chunks at a rate of B_{out}, assuming that $B_{in} > B_{out}$. Note that because of the tit-for-tat policy, the download capacity of a BitTorrent client is mainly dictated by its *outgoing* capacity. In conclusion, the total download capacity will be $S \times B_{out} / N + B_{out}$.

12. **Q:** Give a compelling (technical) argument why the tit-for-tat policy as used in BitTorrent is far from optimal for file sharing in the Internet.

    A: The reasoning is relatively simple. Most BitTorrent clients are operated behind asymmetric links such as provided by ADSL or cable modems. In general, clients are offered a high incoming bandwidth capacity, but no one really expects that clients have services to offer. BitTorrent does not make this assumption, and turns clients into collaborative servers. Having symmetric connections is then a much better match for the tit-for-tat policy.

13. **Q:** We gave two examples of using interceptors in adaptive middleware. What other examples come to mind?

    A: There are several. For example, we could use an interceptor to support mobility. In that case, a request-level interceptor would first look up the current location of a referenced object before the call is forwarded. Likewise, an interceptor can be used to transparently encrypt messages when security is at stake. Another example is when logging is needed. Instead of letting this be handled by the application, we could simply insert a method-specific interceptor that would record specific events before passing a call to the referenced object. More of such example will easily come to mind.

14. **Q:** To what extent are interceptors dependent on the middleware where they are deployed?

    A: In general, interceptors will be highly middleware-dependent. If we consider Fig. 2-0, it is easy to see why: the client stub will most likely be tightly bound to the lower level interfaces offered by the middleware, just as message-level interceptors will be highly dependent on the interaction between middleware and the local operating system. Nevertheless, it is possible to standardize these interfaces, opening the road to developing portable interceptors, albeit often for a specific class of middleware. This last approach has been followed for CORBA.

15. **Q:** Modern cars are stuffed with electronic devices. Give some examples of feed-back control systems in cars.

    A: One obvious one is cruise control. On the one hand this subsystem measures current speed, and when it changes from the required setting, the car is slowed down or speeded up. The anti-lock braking systems (ABS) is another example. By pulsating the brakes of a car, while at the same time regulating the pressure that each wheel is exerting, it is possible to continue steering without losing control because the wheels are blocked. A last example is

formed by the closed circuit of sensors that monitor engine condition. As soon as a dangerous state is reached, a car may come to an automatic halt to prevent the worst.

16. **Q:** Give an example of a self-managing system in which the analysis component is completely distributed or even hidden.

    A: We already came across this type of system: in unstructured peer-to-peer systems where nodes exchange membership information, we saw how a topology could be generated. The analysis component consists of dropping certain links that will not help converge to the intended topology. Similar examples can be found in other such systems as we referred to as well.

17. **Q:** Sketch a solution to automatically determine the best trace length for predicting replication policies in Globule.

    A: An origin server would need to use the traces from T_i to T_{i+1} to check its prediction of policy for that period. It can simply see whether the policy that would have been chosen on the actual access patterns is the same as the one chosen based on the requests in the period T_{i-1} to T_i. This would allow the server to compute the prediction error. By varying the trace length, the origin server would be able find the length for which the prediction is minimal. In this way, we get an automatic determination of the optimal trace length, effectively contributing to the self-managing nature of Globule.

18. **Q:** Using existing software, design and implement a BitTorrent-based system for distributing files to many clients from a single, powerful server. Matters are simplified by using a standard Web server that can operate as tracker.

SOLUTIONS TO CHAPTER 3 PROBLEMS

1. **Q:** In this problem you are to compare reading a file using a single-threaded file server and a multithreaded server. It takes 15 msec to get a request for work, dispatch it, and do the rest of the necessary processing, assuming that the data needed are in a cache in main memory. If a disk operation is needed, as is the case one-third of the time, an additional 75 msec is required, during which time the thread sleeps. How many requests/sec can the server handle if it is single threaded? If it is multithreaded?

   A: In the single-threaded case, the cache hits take 15 msec and cache misses take 90 msec. The weighted average is $2/3 \times 15 + 1/3 \times 90$. Thus the mean request takes 40 msec and the server can do 25 per second. For a multi-threaded server, all the waiting for the disk is overlapped, so every request takes 15 msec, and the server can handle 66 2/3 requests per second.

2. Q: Would it make sense to limit the number of threads in a server process?

A: Yes, for two reasons. First, threads require memory for setting up their own private stack. Consequently, having many threads may consume too much memory for the server to work properly. Another, more serious reason, is that, to an operating system, independent threads tend to operate in a chaotic manner. In a virtual memory system it may be difficult to build a relatively stable working set, resulting in many page faults and thus I/O. Having many threads may thus lead to a performance degradation resulting from page thrashing. Even in those cases where everything fits into memory, we may easily see that memory is accessed following a chaotic pattern rendering caches useless. Again, performance may degrade in comparison to the single-threaded case.

3. Q: In the text, we described a multithreaded file server, showing why it is better than a single-threaded server and a finite-state machine server. Are there any circumstances in which a single-threaded server might be better? Give an example.

A: Yes. If the server is entirely CPU bound, there is no need to have multiple threads. It may just add unnecessary complexity. As an example, consider a telephone directory assistance number for an area with 1 million people. If each (name, telephone number) record is, say, 64 characters, the entire database takes 64 megabytes, and can easily be kept in the server's memory to provide fast lookup.

4. Q: Statically associating only a single thread with a lightweight process is not such a good idea. Why not?

A: Such an association effectively reduces to having only kernel-level threads, implying that much of the performance gain of having threads in the first place, is lost.

5. Q: Having only a single lightweight process per process is also not such a good idea. Why not?

A: In this scheme, we effectively have only user-level threads, meaning that any blocking system call will block the entire process.

6. Q: Describe a simple scheme in which there are as many lightweight processes as there are runnable threads.

A: Start with only a single LWP and let it select a runnable thread. When a runnable thread has been found, the LWP creates another LWP to look for a next thread to execute. If no runnable thread is found, the LWP destroys itself.

7. Q: X designates a user's terminal as hosting the server, while the application is referred to as the client. Does this make sense?

A: Yes, although it may seem a bit confusing. The whole idea is that the server controls the hardware and the application can send requests to manipulate that

hardware. From this perspective the X Window server should indeed reside on the user's machine, having the application acting as its client.

8. **Q:** The X protocol suffers from scalability problems. How can these problems be tackled?

   A: There are essentially two scalability problems. First, numerical scalability is problematic in the sense that too much bandwidth is needed. By using compression techniques, bandwidth can be considerably reduced. Second, there is a geographical scalability problem as an application and the display generally need to synchronize too much. By using cachinng techniques by which effectively state of the display is maintained at the application side, much synchronization traffic can be avoided as the application can inspect the local cache to fins out what the state of the display is.

9. **Q:** Proxies can support replication transparency by invoking each replica, as explained in the text. Can (the server side of) an application be subject to a replicated calls?

   A: Yes: consider a replicated object *A* invoking another (nonreplicated) object *B*. If *A* consists of *k* replicas, an invocation of *B* will be done by each replica. However, *B* should normally be invoked only once. Special measures are needed to handle such replicated invocations.

10. **Q:** Constructing a concurrent server by spawning a process has some advantages and disadvantages compared to multithreaded servers. Mention a few.

    A: An important advantage is that separate processes are protected against each other, which may prove to be necessary as in the case of a superserver handling completely independent services. On the other hand, process spawning is a relatively costly operation that can be saved when using multithreaded servers. Also, if processes do need to communicate, then using threads is much cheaper as in many cases we can avoid having the kernel implement the communication.

11. **Q:** Sketch the design of a multithreaded server that supports multiple protocols using sockets as its transport-level interface to the underlying operating system.

    A: A relatively simple design is to have a single thread *T* waiting for incoming transport messages (TPDUs). If we assume the header of each TPDU contains a number identifying the higher-level protocol, the tread can take the payload and pass it to the module for that protocol. Each such module has a separate thread waiting for this payload, which it treats as an incoming request. After handling the request, a response message is passed to *T*, which, in turn, wraps it in a transport-level message and sends it to tthe proper destination.

12. **Q:** How can we prevent an application from circumventing a window manager, and thus being able to completely mess up a screen?

    A: Use a microkernel approach by which the windowing system including the window manager are run in such a way that all window operations are required to go through the kernel. In effect, this is the essence of transferring the client-server model to a single computer.

13. **Q:** Is a server that maintains a TCP/IP connection to a client stateful or stateless?

    A: Assuming the server maintains no other information on that client, one could justifiably argue that the server is stateless. The issue is that not the server, but the transport layer at the server maintains state on the client. What the local operating systems keep track of is, in principle, of no concern to the server.

14. **Q:** Imagine a Web server that maintains a table in which client IP addresses are mapped to the most recently accessed Web pages. When a client connects to the server, the server looks up the client in its table, and if found, returns the registered page. Is this server stateful or stateless?

    A: It can be strongly argued that this is a stateless server. The important issue with stateless designs is not if *any* information is maintained by the server on its clients, but instead whether that information is needed for correctness. In this example, if the table is lost for what ever reason, the client and server can still properly interact as if nothing happened. In a stateful design, such an interaction would be possible only after the server had recovered from a possible fault.

15. **Q:** Strong mobility in UNIX systems could be supported by allowing a process to fork a child on a remote machine. Explain how this would work.

    A: Forking in UNIX means that a complete image of the parent is copied to the child, meaning that the child continues just after the call to fork. A similar approach could be used for remote cloning, provided the target platform is exactly the same as where the parent is executing. The first step is to have the target operating system reserve resources and create the appropriate process and memory map for the new child process. After this is done, the parent's image (in memory) can be copied, and the child can be activated. (It should be clear that we are ignoring several important details here.)

16. **Q:** In Fig. 3-0 it is suggested that strong mobility cannot be combined with executing migrated code in a target process. Give a counterexample.

    A: If strong mobility takes place through thread migration, it should be possible to have a migrated thread be executed in the context of the target process.

17. **Q:** Consider a process P that requires access to file F which is locally available on the machine where P is currently running. When P moves to another machine, it still requires access to F. If the file-to-machine binding is fixed, how could the systemwide reference to F be implemented?

A: The simplest solution is to create a separate process Q that handles remote requests for F. Process P is offered the same interface to F as before, for example in the form of a proxy. Effectively, process Q operates as a file server.

18. **Q:** Describe in detail how TCP packets flow in the case of TCP handoff, along with the information on source and destination addresses in the various headers.

A: There are various ways of doing this, but a simple one is to let a front end to execute a three-way handshake and from there on forward packets to a selected server. That server sends TCP PDUs in which the source address corresponds to that of the front end. An alternative is to forward the first packet to a server. Note, however, that in this case the front end will continue to stay in the loop. The advantage of this scheme is that the selected server builds up the required TCP state (such as the sequence numbers to be used) instead of obtaining this information from the front end as in the first scenario.

SOLUTIONS TO CHAPTER 4 PROBLEMS

1. **Q:** In many layered protocols, each layer has its own header. Surely it would be more efficient to have a single header at the front of each message with all the control in it than all these separate headers. Why is this not done?

A: Each layer must be independent of the other ones. The data passed from layer $k+1$ down to layer k contains both header and data, but layer k cannot tell which is which. Having a single big header that all the layers could read and write would destroy this transparency and make changes in the protocol of one layer visible to other layers. This is undesirable.

2. **Q:** Why are transport-level communication services often inappropriate for building distributed applications?

A: They hardly offer distribution transparency meaning that application developers are required to pay significant attention to implementing communication, often leading to proprietary solutions. The effect is that distributed applications, for example, built directly on top of sockets are difficult to port and to interoperate with other applications.

3. **Q:** A reliable multicast service allows a sender to reliably pass messages to a collection of receivers. Does such a service belong to a middleware layer, or should it be part of a lower-level layer?

A: In principle, a reliable multicast service could easily be part of the transport

layer, or even the network layer. As an example, the unreliable IP multicasting service is implemented in the network layer. However, because such services are currently not readily available, they are generally implemented using transport-level services, which automatically places them in the middleware. However, when taking scalability into account, it turns out that reliability can be guaranteed only if application requirements are considered. This is a strong argument for implementing such services at higher, less general layers.

4. Q: Consider a procedure *incr* with two integer parameters. The procedure adds one to each parameter. Now suppose that it is called with the same variable twice, for example, as *incr*(*i*, *i*). If *i* is initially 0, what value will it have afterward if call-by-reference is used? How about if copy/restore is used?

A: If call by reference is used, a pointer to *i* is passed to *incr*. It will be incremented two times, so the final result will be two. However, with copy/restore, *i* will be passed by value twice, each value initially 0. Both will be incremented, so both will now be 1. Now both will be copied back, with the second copy overwriting the first one. The final value will be 1, not 2.

5. Q: C has a construction called a union, in which a field of a record (called a struct in C) can hold any one of several alternatives. At run time, there is no sure-fire way to tell which one is in there. Does this feature of C have any implications for remote procedure call? Explain your answer.

A: If the runtime system cannot tell what type value is in the field, it cannot marshal it correctly. Thus unions cannot be tolerated in an RPC system unless there is a tag field that unambiguously tells what the variant field holds. The tag field must not be under user control.

6. Q: One way to handle parameter conversion in RPC systems is to have each machine send parameters in its native representation, with the other one doing the translation, if need be. The native system could be indicated by a code in the first byte. However, since locating the first byte in the first word is precisely the problem, can this actually work?

A: First of all, when one computer sends byte 0, it always arrives in byte 0. Thus the destination computer can simply access byte 0 (using a byte instruction) and the code will be in it. It does not matter whether this is the low-order byte or the high-order byte. An alternative scheme is to put the code in all the bytes of the first word. Then no matter which byte is examined, the code will be there.

7. Q: Assume a client calls an asynchronous RPC to a server, and subsequently waits until the server returns a result using another asynchronous RPC. Is this approach the same as letting the client execute a normal RPC? What if we replace the asynchronous RPCs with asynchronous RPCs?

A: No, this is not the same. An asynchronous RPC returns an acknowledgment

to the caller, meaning that after the first call by the client, an additional message is sent across the network. Likewise, the server is acknowledged that its response has been delivered to the client. Two asynchronous RPCs may be the same, provided reliable communication is guaranteed. This is generally not the case.

8. **Q:** Instead of letting a server register itself with a daemon as in DCE, we could also choose to always assign it the same endpoint. That endpoint can then be used in references to objects in the server's address space. What is the main drawback of this scheme?

   A: The main drawback is that it becomes much harder to dynamically allocate objects to servers. In addition, many endpoints need to be fixed, instead of just one (i.e., the one for the daemon). For machines possibly having a large number of servers, static assignment of endpoints is not a good idea.

9. **Q:** Would it be useful to also make a distinction between static and dynamic RPCs?

   A: Yes, for the same reason it is useful with remote object invocations: it simply introduces more flexibility. The drawback, however, is that much of the distribution transparency is lost for which RPCs were introduced in the first place.

10. **Q:** Describe how connectionless communication between a client and a server proceeds when using sockets.

    A: Both the client and the server create a socket, but only the server binds the socket to a local endpoint. The server can then subsequently do a blocking read call in which it waits for incoming data from any client. Likewise, after creating the socket, the client simply does a blocking call to write data to the server. There is no need to close a connection.

11. **Q:** Explain the difference between the primitives mpi_bsend and mpi_isend in MPI.

    A: The primitive mpi_bsend uses buffered communication by which the caller passes an entire buffer containing the messages to be sent, to the local MPI runtime system. When the call completes, the messages have either been transferred, or copied to a local buffer. In contrast, with mpi_isend, the caller passes only a pointer to the message to the local MPI runtime system after which it immediately continues. The caller is responsible for not overwriting the message that is pointed to until it has been copied or transferred.

12. **Q:** Suppose that you could make use of only transient asynchronous communication primitives, including only an asynchronous receive primitive. How would you implement primitives for transient *synchronous* communication?

    A: Consider a synchronous send primitive. A simple implementation is to send a message to the server using asynchronous communication, and subsequently

let the caller continuously poll for an incoming acknowledgment or response from the server. If we assume that the local operating system stores incoming messages into a local buffer, then an alternative implementation is to block the caller until it receives a signal from the operating system that a message has arrived, after which the caller does an asynchronous receive.

13. **Q:** Suppose that you could make use of only transient synchronous communication primitives. How would you implement primitives for transient *asynchronous* communication?

A: This situation is actually simpler. An asynchronous send is implemented by having a caller append its message to a buffer that is shared with a process that handles the actual message transfer. Each time a client appends a message to the buffer, it wakes up the send process, which subsequently removes the message from the buffer and sends it its destination using a blocking call to the original send primitive. The receiver is implemented in a similar fashion by offering a buffer that can be checked for incoming messages by an application.

14. **Q:** Does it make sense to implement persistent asynchronous communication by means of RPCs?

A: Yes, but only on a hop-to-hop basis in which a process managing a queue passes a message to a next queue manager by means of an RPC. Effectively, the service offered by a queue manager to another is the storage of a message. The calling queue manager is offered a proxy implementation of the interface to the remote queue, possibly receiving a status indicating the success or failure of each operation. In this way, even queue managers see only queues and no further communication.

15. **Q:** In the text we stated that in order to automatically start a process to fetch messages from an input queue, a daemon is often used that monitors the input queue. Give an alternative implementation that does not make use of a daemon.

A: A simple scheme is to let a process on the receiver side check for any incoming messages each time that process puts a message in its own queue.

16. **Q:** Routing tables in IBM WebSphere, and in many other message-queuing systems, are configured manually. Describe a simple way to do this automatically.

A: The simplest implementation is to have a centralized component in which the topology of the queuing network is maintained. That component simply calculates all best routes between pairs of queue managers using a known routing algorithm, and subsequently generates routing tables for each queue manager. These tables can be downloaded by each manager separately. This approach works in queuing networks where there are only relatively few, but possibly widely dispersed, queue managers.

A more sophisticated approach is to decentralize the routing algorithm, by having each queue manager discover the network topology, and calculate its own best routes to other managers. Such solutions are widely applied in computer networks. There is no principle objection for applying them to message-queuing networks.

17. **Q:** With persistent communication, a receiver generally has its own local buffer where messages can be stored when the receiver is not executing. To create such a buffer, we may need to specify its size. Give an argument why this is preferable, as well as one against specification of the size.

   A: Having the user specify the size makes its implementation easier. The system creates a buffer of the specified size and is done. Buffer management becomes easy. However, if the buffer fills up, messages may be lost. The alternative is to have the communication system manage buffer size, starting with some default size, but then growing (or shrinking) buffers as need be. This method reduces the chance of having to discard messages for lack of room, but requires much more work of the system.

18. **Q:** Explain why transient synchronous communication has inherent scalability problems, and how these could be solved.

   A: The problem is the limited geographical scalability. Because synchronous communication requires that the caller is blocked until its message is received, it may take a long time before a caller can continue when the receiver is far away. The only way to solve this problem is to design the calling application so that it has other useful work to do while communication takes place, effectively establishing a form of asynchronous communication.

19. **Q:** Give an example where multicasting is also useful for discrete data streams.

   A: Passing a large file to many users as is the case, for example, when updating mirror sites for Web services or software distributions.

20. **Q:** Suppose that in a sensor network measured temperatures are not timestamped by the sensor, but are immediately sent to the operator. Would it be enough to guarantee only a maximum end-to-end delay?

   A: Not really if we assume that the operator would still need to know when the measurement took place. In this case, a timestamp can be attached when the measurement is received, but this would mean that we should also have guarantees for minimum end-to-end delays.

21. **Q:** How could you guarantee a maximum end-to-end delay when a collection of computers is organized in a (logical or physical) ring?

   A: We let a token circulate the ring. Each computer is permitted to send data across the ring (in the same direction as the token) only when holding the token. Moreover, no computer is allowed to hold the token for more than T

seconds. Effectively, if we assume that communication between two adjacent computers is bounded, then the token will have a maximum circulation time, which corresponds to a maximum end-to-end delay for each packet sent.

22. **Q:** How could you guarantee a minimum end-to-end delay when a collection of computers is organized in a (logical or physical) ring?

    A: Strangely enough, this is much harder than guaranteeing a maximum delay. The problem is that the receiving computer should, in principle, not receive data before some elapsed time. The only solution is to buffer packets as long as necessary. Buffering can take place either at the sender, the receiver, or somewhere in between, for example, at intermediate stations. The best place to temporarily buffer data is at the receiver, because at that point there are no more unforeseen obstacles that may delay data delivery. The receiver need merely remove data from its buffer and pass it to the application using a simple timing mechanism. The drawback is that enough buffering capacity needs to be provided.

23. **Q:** Despite that multicasting is technically feasible, there is very little support to deploy it in the Internet. The answer to this problem is to be sought in down-to-earth business models: no one really knows how to make money out of multicasting. What scheme can you invent?

    A: The problem is mainly caused by ISPs, as they see no reason to save on bandwidth (their clients are paying anyway). However, matters may change in scenarios such as the following. An Internet broadcasting service pays for a certain quality-of-service as promised by various ISPs. Each of these ISPs will see a drop in their income when they cannot meet these QoS requirements. At this point, they may now have an incentive to start deploying multicasting as they can offer better (and guaranteed) service.

24. **Q:** Normally, application-level multicast trees are optimized with respect stretch, which is measured in terms of delay or hop counts. Give an example where this metric could lead to very poor trees.

    A: The underlying assumption with stretch is that communication delays predominate performance. However, in the case of, for example, video broadcasting, it is the available which counts. In that case, we would like to construct trees that *maximize* costs (measured in terms of bandwidth).

25. **Q:** When searching for files in an unstructured peer-to-peer system, it may help to restrict the search to nodes that have similar files as yourself. Explain how gossiping can help to find those nodes.

    A: The idea is very simple: if, during gossiping, nodes exchange membership information, every node will eventually get to know about all other nodes in the system. Each time it discovers a new node, it can be evaluated with respect to its semantic proximity, for example, by counting the number of files in

common. The semantically nearest nodes are then selected for submitting a search query.

SOLUTIONS TO CHAPTER 5 PROBLEMS

1. **Q:** Give an example of where an address of an entity E needs to be further resolved into another address to actually access E.

   A: IP addresses in the Internet are used to address hosts. However, to access a host, its IP address needs to be resolved to, for example, an Ethernet address.

2. **Q:** Would you consider a URL such as *http://www.acme.org/index.html* to be location independent? What about *http://www.acme.nl/index.html*?

   A: Both names can be location independent, although the first one gives fewer hints on the location of the named entity. Location independent means that the name of the entity is independent of its address. By just considering a name, nothing can be said about the address of the associated entity.

3. **Q:** Give some examples of true identifiers.

   A: Examples are ISBN numbers for books, identification numbers for software and hardware products, employee numbers within a single organization, and Ethernet addresses (although some addresses are used to identify a machine instead of just the Ethernet board).

4. **Q:** Is an identifier allowed to contain information on the entity it refers to?

   A: Yes, but that information is not allowed to change, because that would imply changing the identifier. The old identifier should remain valid, so that changing it would imply that an entity has two identifiers, violating the second property of identifiers.

5. **Q:** Outline an efficient implementation of globally unique identifiers.

   A: Such identifiers can be generated locally in the following way. Take the network address of the machine where the identifier is generated, append the local time to that address, along with a generated pseudo-random number. Although, in theory, it is possible that another machine in the world can generate the same number, chances that this happens are negligible.

6. **Q:** Consider the Chord system as shown in Fig. 5-0 and assume that node 7 has just joined the network. What would its finger table be and would there be any changes to other finger tables?

   A: Let us first consider the finger table for node 7. Using the same method as we introduced when discussing Chord, it can be seen that this table will be [9, 9, 11, 18, 28]. For example, entry #2 is computed as $succ(7 + 2 1) = succ(9) = 9$. More tables will need to change, however, in

particular those of node 4 (which becomes [7,7,9,14,28]), node 21 ([28,28,28,1,7]) and node 1 ([4,4,7,9,18]).

7. **Q:** Consider a Chord DHT-based system for which k bits of an m-bit identifier space have been reserved for assigning to superpeers. If identifiers are randomly assigned, how many superpeers can one expect to have in an N-node system?

   A: The total number of superpeers in an overlay of N nodes will be equal to $\min\{2^{k-m}N, 2^k\}$.

8. **Q:** If we insert a node into a Chord system, do we need to instantly update all the finger tables?

   A: Fortunately not. Consider the previous question and assume that only the finger table of node #7 is installed and that the rest are kept as they are. The worst that can happen is that a request to look up, say, key 5, is routed to node #9 instead of #7. However, node #9 knows that node #7 has joined the system and can therefore take corrective action.

9. **Q:** What is a major drawback of recursive lookups when resolving a key in a DHT-based system?

   A: A problem is that the requesting client will never be able to discover what went wrong when no answer is returned. Somewhere along the route that corresponds to resolving the key, a message may have been lost or a node may have failed. For this reason, an iterative lookup is sometimes preferred: the client will know exactly which part of the lookup did not come through and may be able to choose an alternative node to help out.

10. **Q:** A special form of locating an entity is called anycasting, by which a service is identified by means of an IP address (see, for example, Sending a request to an anycast address, returns a response from a server implementing the service identified by that anycast address. Outline the implementation of an anycast service based on the hierarchical location service described in Sec. 5.2.4.

    A: Each service has a unique identifier associated with it. Any server implementing that service, inserts its network-level address into the location service at the directory node of the leaf domain in which the server resides. Lookup requests use the identifier of the service, and will automatically return the nearest server implementing that service.

11. **Q:** Considering that a two-tiered home-based approach is a specialization of a hierarchical location service, where is the root?

    A: The root is formed jointly by all home locations, but is partitioned in such a way that each mobile entity has its own root server.

12. **Q:** Suppose that it is known that a specific mobile entity will almost never move outside domain D, and if it does, it can be expected to return soon. How can this information be used to speed up the lookup operation in a hierarchical location service?

    A: Simply encode the domain D in the identifier for the entity that is used for the lookup operation. The operation can then be immediately forwarded to the directory node *dir(D)*, from where the search continues.

13. **Q:** In a hierarchical location service with a depth of k, how many location records need to be updated at most when a mobile entity changes its location?

    A: Changing location can be described as the combination of an insert and a delete operation. An insert operation requires that at worst $k+1$ location records are to be changed. Likewise, a delete operation also requires changing $k+1$ records, where the record in the root is shared between the two operations. This leads to a total of $2k+1$ records.

14. **Q:** Consider an entity moving from location A to B, while passing several intermediate locations where it will reside for only a relatively short time. When arriving at B, it settles down for a while. Changing an address in a hierarchical location service may still take a relatively long time to complete, and should therefore be avoided when visiting an intermediate location. How can the entity be located at an intermediate location?

    A: Combine the hierarchical location service with forwarding pointers. When the entity starts to move, it leaves behind a forwarding pointer at A to its next (intermediate) location. Each time it moves again, a forwarding pointer is left behind. Upon arrival in B, the entity inserts its new address into the hierarchical location service. The chain of pointers is subsequently cleaned up, and the address in A is deleted.

15. **Q:** The root node in hierarchical location services may become a potential bottleneck. How can this problem be effectively circumvented?

    A: An important observation is that we are using only random bit strings as identifiers. As a result, we can easily partition the identifier space and install a separate root node for each part. In addition, the partitioned root node should be spread across the network so that accesses to it will also be spread.

16. **Q:** Give an example of how the closure mechanism for a URL could work.

    A: Assuming a process knows it is dealing with a URL, it first extracts the scheme identifier from the URL, such as the string *ftp:*. It can subsequently look up this string in a table to find an interface to a local implementation of the FTP protocol. The next part of the closure mechanism consists of extracting the host name from the URL, such as *www.cs.vu.nl*, and passing that to the local DNS name server. Knowing where and how to contact the DNS server is an important part of the closure mechanism. It is often hard-coded into the

URL name resolver that the process is executing. Finally, the last part of the URL, which refers to a file that is to be looked up, is passed to the identified host. The latter uses its own local closure mechanism to start the name resolution of the file name.

17. Q: Explain the difference between a hard link and a soft link in UNIX systems. Are there things that can be done with a hard link that cannot be done with a soft link or vice versa?

A: A hard link is a named entry in a directory file pointing to the same file descriptor as another named entry (in possibly a different directory). A symbolic link is a file containing the (character string) name of another file. With a soft link you can link to a different disk partition or even to a different machine.

18. Q: High-level name servers in DNS, that is, name servers implementing nodes in the DNS name space that are close to the root, generally do not support recursive name resolution. Can we expect much performance improvement if they did?

A: Probably not: because the high-level name servers constitute the global layer of the DNS name space, it can be expected that changes to that part of the name space do not occur often. Consequently, caching will be highly effective, and much long-haul communication will be avoided anyway. Note that recursive name resolution for low-level name servers is important, because in that case, name resolution can be kept local at the lower-level domain in which the resolution is taking place.

19. Q: Explain how DNS can be used to implement a home-based approach to locating mobile hosts.

A: The DNS name of a mobile host would be used as (rather poor) identifier for that host. Each time the name is resolved, it should return the current IP address of the host. This implies that the DNS server responsible for providing that IP address will act as the host's name server. Each time the host moves, it contacts this home server and provides it with its current address. Note that a mechanism should be available to avoid caching of the address. In other words, other name servers should be told *not* to cache the address found.

20. Q: How is a mounting point looked up in most UNIX systems?

A: By means of a mount table that contains an entry pointing to the mount point. This means that when a mounting point is to be looked up, we need to go through the mount table to see which entry matches a given mount point.

21. Q: Consider a distributed file system that uses per-user name spaces. In other words, each user has his own, private name space. Can names from such name spaces be used to share resources between two different users?

A: Yes, provided names in the per-user name spaces can be resolved to names

in a shared, global name space. For example, two identical names in different name spaces are, in principle, completely independent and may refer to different entities. To share entities, it is necessary to refer to them by names from a shared name space. For example, Jade relies on DNS names and IP addresses that can be used to refer to shared entities such as FTP sites.

22. **Q:** Consider DNS. To refer to a node *N* in a subdomain implemented as a different zone than the current domain, a name server for that zone needs to be specified. Is it always necessary to include a resource record for that server's address, or is it sometimes sufficient to provide only its domain name?

   A: When the name server is represented by a node *NS* in a domain other than the one in which *N* is contained, it is enough to give only its domain name. In that case, the name can be looked up by a separate DNS query. This is not possible when *NS* lies in the same subdomain as *N*, for in that case, you would need to contact the name server to find out its address.

23. **Q:** Counting common files is a rather naive way of defining semantic proximity. Assume you were to build semantic overlay networks based on text documents, what other semantic proximity function can you think of?

   A: One intriguing one is to have a look at actual content when possible. In the case of documents, one could look at similarity functions derived from information retrieval, such as the Vector Space Model (VSM).

24. **Q:** Set up your own DNS server. Install BIND on either a Windows or UNIX machine and configure it for a few simple names. Test your configuration using tools such as the Domain Information Groper (DIG). Make sure your DNS database includes records for name servers, mail servers, and standard servers. Note that if you are running BIND on a machine with host name *HOSTNAME*, you should be able to resolve names of the form *RESOURCE-NAME.HOSTNAME*.

SOLUTIONS TO CHAPTER 6 PROBLEMS

1. **Q:** Name at least three sources of delay that can be introduced between WWV broadcasting the time and the processors in a distributed system setting their internal clocks.

   A: First we have signal propagation delay in the atmosphere. Second we might have collision delay while the machines with the WWV receivers fight to get on the Ethernet. Third, there is packet propagation delay on the LAN. Fourth, there is delay in each processor after the packet arrives, due to interrupt processing and internal queueing delays.

2. **Q:** Consider the behavior of two machines in a distributed system. Both have clocks that are supposed to tick 1000 times per millisecond. One of them actually does, but the other ticks only 990 times per millisecond. If UTC updates come in once a minute, what is the maximum clock skew that will occur?

A: The second clock ticks 990,000 times per second, giving an error of 10 msec per second. In a minute this error has grown to 600 msec. Another way of looking at it is that the second clock is one percent slow, so after a minute it is off by 0.01×60 sec, or 600 msec.

3. **Q:** One of the modern devices that have (silently) crept into distributed systems are GPS receivers. Give examples of distributed applications that can make use of GPS information.

A: One typical example that comes to mind is sports and health care. There are now GPS-based body-area networks that allow a person to keep track of his pace while exercising an outdoors sports. These networks are often augmented with heart rate monitors and can be hooked up to a computer to downlaod the sensored data for further analysis. Another example is formed by car-navigation equipment, which is generally based on GPS receivers. Hooked up to a (portable) personal computer that can be connected to the Internet, maps and such can be continously updated. Related are GPS-based distributed systems for tracking the movement of cars and trucks.

4. **Q:** When a node synchronizes its clock to that of another node, it is generally a good idea to take previous measurements into account as well. Why? Also, give an example of how such past readings could be taken into account.

A: The obvious reason is that there may be an error in the current reading. Assuming that clocks need only be gradually adjusted, one possibility is to consider the last N values and compute a median or average. If the measured value falls outside a current interval, it is not taken into account (but is added to the list). Likewise, a new value can be computed by taking a weighted average, or an aging algorithm.

5. **Q:** Add a new message to Fig. 6-0 that is concurrent with message A, that is, it neither happens before A nor happens after A.

A: The solution cannot involve 0 or it would be ordered. Thus it must be a message from 1 to 2 or from 2 to 1. If it departs or arrives from 1 after 16, it will be ordered with respect to A, so it must depart or arrive before 16. The possibilities are a message leaving process 2 at 0 and arriving at process 1 at 8, or a message leaving process 1 at 0 and arriving at process 2 at 10. Both of these are concurrent with A.

6. Q: To achieve totally-ordered multicasting with Lamport timestamps, is it strictly necessary that each message is acknowledged?

A: No, it is sufficient to multicast any other type of message, as long as that message has a timestamp larger than the received message. The condition for delivering a message m to the application, is that another message has been received from each other process with a large timestamp. This guarantees that there are no more messages underway with a lower timestamp.

7. Q: Consider a communication layer in which messages are delivered only in the order that they were sent. Give an example in which even this ordering is unnecessarily restrictive.

A: Imagine the transfer of a large image which, to that end, has been divided into consecutive blocks. Each block is identified by its position in the original image, and possibly also its width and height. In that case, FIFO ordering is not necessary, as the receiver can simply paste each incoming block into the correct position.

8. Q: Many distributed algorithms require the use of a coordinating process. To what extent can such algorithms actually be considered distributed? Discuss.

A: In a centralized algorithm, there is often one, fixed process that acts as coordinator. Distribution comes from the fact that the other processes run on different machines. In distributed algorithms with a nonfixed coordinator, the coordinator is chosen (in a distributed fashion) among the processes that form part of the algorithm. The fact that there is a coordinator does not make the algorithm less distributed.

9. Q: In the centralized approach to mutual exclusion (Fig. 6-0), upon receiving a message from a process releasing its exclusive access to the resources it was using, the coordinator normally grants permission to the first process on the queue. Give another possible algorithm for the coordinator.

A: Requests could be associated with priority levels, depending on their importance. The coordinator could then grant the highest priority request first.

10. Q: Consider Fig. 6-0 again. Suppose that the coordinator crashes. Does this always bring the system down? If not, under what circumstances does this happen? Is there any way to avoid the problem and make the system able to tolerate coordinator crashes?

A: Suppose that the algorithm is such that every request is answered immediately, either with permission or with denial. If there are no processes accessing resources and no processes queued, then a crash is not fatal. The next process to request permission will fail to get any reply at all, and can initiate the election of a new coordinator. The system can be made even more robust by having the coordinator store every incoming request on disk *before* sending back a reply. In this way, in the event of a crash, the new coordinator can reconstruct

the list of accessed resources and the queue by reading the file from the disk.

11. **Q:** Ricart and Agrawala's algorithm has the problem that if a process has crashed and does not reply to a request from another process to access a resources, the lack of response will be interpreted as denial of permission. We suggested that all requests be answered immediately to make it easy to detect crashed processes. Are there any circumstances where even this method is insufficient? Discuss.

    A: Suppose that a process denies permission and then crashes. The requesting process thinks that it is alive, but permission will never come. One way out is to have the requester not actually block, but rather go to sleep for a fixed period of time, after which it polls all processes that have denied permission to see if they are still running.

12. **Q:** How do the entries in Fig. 6-0 change if we assume that the algorithms can be implemented on a LAN that supports hardware broadcasts?

    A: Only the entries for the distributed case change. Because sending a point-to-point message is as expensive as doing a broadcast, we need only send one broadcast message to all processes requesting access to the resource. Likewise, only one release broadcast message is needed. The delay becomes $1 + (n - 1)$: one delay coming from the broadcast request, and an additional $n - 1$ as we still need to receive a message from each other process before being allowed to access the resource.

13. **Q:** A distributed system may have multiple, independent resources. Imagine that process 0 wants to access resource A and process 1 wants to access resource B. Can Ricart and Agrawala's algorithm lead to deadlocks? Explain your answer.

    A: It depends on the ground rules. If processes access resources strictly sequentially, that is, a process holding a resource may not attempt to access another one, then there is no way that it can block while holding a resource that some other process wants. The system is then deadlock free. On the other hand, if process 0 may hold resource A and then try to access resource B, a deadlock can occur if some other process tries to acquire them in the reverse order. The Ricart and Agrawala algorithm itself does not contribute to deadlock since each resource is handled independently of all the others.

14. **Q:** Suppose that two processes detect the demise of the coordinator simultaneously and both decide to hold an election using the bully algorithm. What happens?

    A: Each of the higher-numbered processes will get two *ELECTION* messages, but will ignore the second one. The election will proceed as usual.

15. **Q:** In Fig. 6-0 we have two *ELECTION* messages circulating simultaneously. While it does no harm to have two of them, it would be more elegant if one could be killed off. Devise an algorithm for doing this without affecting the operation of the basic election algorithm.

    A: When a process receives an *ELECTION* message, it checks to see who started it. If it itself started it (i.e., its number is at the head of the list), it turns the message into a *COORDINATOR* message as described in the text. If it did not start any *ELECTION* message, it adds its process number and forwards it along the ring. However, if it did send its own *ELECTION* message earlier and it has just discovered a competitor, it compares the originator's process number with its own. If the other process has a lower number, it discards the message instead of passing it on. If the competitor is higher, the message is forwarded in the usual way. In this way, if multiple *ELECTION* messages are started, the one whose first entry is highest survives. The rest are killed off along the route.

16. **Q:** UNIX systems provide many facilities to keep computers in synch. Notably the combination of the *crontab* tool (which allows to automatically schedule operations) and various synchronization commands are powerful. Configure a UNIX that keeps the local time accurate with in the range of a single second. Likewise, configure an automatic backup facility by which a number of crucial files are automatically transferred to a remote machine once every 5 minutes. Your solution should be efficient when it comes to bandwidth usage.

SOLUTIONS TO CHAPTER 7 PROBLEMS

1. **Q:** Access to shared Java objects can be serialized by declaring its methods as being synchronized. Is this enough to guarantee serialization when such an object is replicated?

   A: No. The problem is that access to each replica is serialized. However, different operations at different replicas may be executed at the same time, leaving the replicated instance variables in an inconsistent state.

2. **Q:** Explain in your own words what the main reason is for actually considering weak consistency models.

   A: Weak consistency models come from the need to replicate for performance. However, efficient replication can be done only if we can avoid global synchronizations, which, in turn, can be achieved by loosening consistency constraints.

3. **Q:** Explain how replication in DNS takes place, and why it actually works so well.

   A: The basic idea is that name servers cache previously looked up results.

These results can be kept in a cache for a long time, because DNS makes the assumption that name-to-address mappings do not change often.

4. **Q:** During the discussion of consistency models, we often referred to the contract between the software and data store. Why is such a contract needed?

   A: If a program expects a sequentially consistent data store and cannot live with anything less, the store must provide sequential consistency. However, to improve performance, some systems provide a weaker model. It is then essential that the software agrees to abide by the rules imposed by this model. Generally, it means that programs obeying the rules will perceive what looks like a sequentially consistent data store.

5. **Q:** Given the replicas in Fig. 7-0, what would need to be done to finalize the values in the conit such that both *A* and *B* see the same result?

   A: In this case it is relatively simple: if *A* and *B* exchange their list of tentative operations and subsequently order them according the time, then both would get to see the same result.

6. **Q:** In Fig. 7-0, is 001110 a legal output for a sequentially consistent memory? Explain your answer.

   A: Yes. If the processes run in the order (a), (c), (b), this result is obtained.

7. **Q:** It is often argued that weak consistency models impose an extra burden for programmers. To what extent is this statement actually true?

   A: It really depends. Many programmers are used to protect their shared data through synchronization mechanisms such as locks or transactions. The main idea is that they require a coarser grain of concurrency than one offered at the level of only read and write operations. However, programmers do expect that operations on synchronization variables adhere to sequential consistency.

8. **Q:** Does totally ordered multicasting by means of a sequencer and for the sake of consistency in active replication, violate the end-to-end argument in system design?

   A: Yes. The end-to-end argument states that problems should be solved at the same level in which they occur. In this case, we are dealing with the problem of totally ordered multicasting for achieving consistency in active replication. In primary-based protocols, consistency is achieved by first forwarding all operations to the primary. Using a sequencer, we are actually doing the same but at a lower level of abstraction. In this case, it may have been better to use a primary-based protocol in which updates are propagated by sending operations.

9. **Q:** What kind of consistency would you use to implement an electronic stock market? Explain your answer.

   A: Causal consistency is probably enough. The issue is that reactions to

changes in stock values should be consistent. Changes in stocks that are independent can be seen in different orders.

10. **Q:** Consider a personal mailbox for a mobile user, implemented as part of a wide-area distributed database. What kind of client-centric consistency would be most appropriate?

    A: All of them, actually. What it boils down to is that the owner should always see the same mailbox, no matter whether he is reading or updating it. In fact, the simplest implementation for such a mailbox may well be that of a primary-based local-write protocol, where the primary is always located on the user's mobile computer.

11. **Q:** Describe a simple implementation of read-your-writes consistency for displaying Web pages that have just been updated.

    A: The simplest implementation is to let the browser always check whether it is displaying the most recent version of a page. This requires sending a request to the Web server. This scheme is simple as it is already implemented by many systems.

12. **Q:** To make matters simple, we assumed that there were no write-write conflicts in Bayou. Of course, this is an unrealistic assumption. Explain how conflicts may happen.

    A: There are many occasions in which conflicts can occur. For example, nothing prevents a client of using a shared agenda. In that case, updates may consist of scheduling a meeting at a time where the client has already scheduled something else, but this information had not yet been propagated to other replicas.

13. **Q:** When using a lease, is it necessary that the clocks of a client and the server, respectively, are tightly synchronized?

    A: No. If the client takes a pessimistic view concerning the level at which its clock is synchronized with that of the server, it will attempt to obtain a new lease far before the current one expires.

14. **Q:** We have stated that totally ordered multicasting using Lamport's logical clocks does not scale. Explain why.

    A: Lamport's way totally ordered multicasting requires that all servers are up and running, effectively hindering performance when one of them turns out to be slow or has crashed. This will have to be detected by all other servers. As the number of servers grows, this problem is aggravated.

15. **Q:** Show that, in the case of continuous consistency, having a server S_k advance its view $TW_k(i, k)$ whenever it receives a fresh update that would increase $TW(k, k) - TW_k(i, k)$ beyond $\delta_i / N - 1)$, ensures that $v(t) - v_i \leq \delta_i$.

    A: That this advancement indeed yields correctness can easily be seen as

follows:

$$v(t) - v_i = (v(0) + \sum_{k=1}^{N} TW[k, k]) + (v(0) + \sum_{k=1}^{N} TW[i, k])$$

$$= \sum_{k=1}^{N} (TW[k, k] - TW[i, k])$$

$$\leq \sum_{k=1}^{N} [TW[k, k] - TW_k[i, k]] \leq (N - 1) \times \delta_i / (N - 1) = \delta_i$$

Note that the factor $(N - 1)$ comes from the fact that $TW[k, k] - TW_k[k, k] = 0$, by which we can eliminate a term in the summation.

16. **Q:** For continuous consistency, we have assumed that each write only increases the value of data item x. Sketch a solution in which it is also possible to decrease x's value.

    A: The situation is relatively simple if we separate positive-valued updates from negative-valued ones and keep separate accounts for each of them. In particular, we keep track of:

    $$TWN[i, j] = \sum \{weight(W) \,|\, weight(W) < 0 \,\&\, origin(W) = S_j \,\&\, W \in L_i\}$$

    $$TWP[i, j] = \sum \{weight(W) \,|\, weight(W) > 0 \,\&\, origin(W) = S_j \,\&\, W \in L_i\}$$

    Note that $TWP[i, j] \equiv TW[i, j]$. Again, each node keeps track of views $TWP_k[i, k]$ and $TWN_k[i, k]$, respectively, and advances its view when it notices that a fresh write would either increase $|TWN[k, k] - TWN_k[i, k]|$ or $|TWP[k, k] - TWP_k[i, k]|$ beyond $\delta_i / (N - 1)$.

17. **Q:** Consider a nonblocking primary-backup protocol used to guarantee sequential consistency in a distributed data store. Does such a data store always provide read-your-writes consistency?

    A: No. As soon as the updating process receives an acknowledgment that its update is being processed, it may disconnect from the data store and reconnect to another replica. No guarantees are given that the update has already reached that replica. In contrast, with a blocking protocol, the updating process can disconnect only after its update has been fully propagated to the other replicas as well.

18. **Q:** For active replication to work in general, it is necessary that all operations be carried out in the same order at each replica. Is this ordering always necessary?

    A: No. Consider read operations that operate on nonmodified data or commutative write operations. In principle, such situations allow ordering to be different at different replicas. However, it can be hard or impossible to detect

whether, for example, two write operations are commutative.

19. **Q:** To implement totally ordered multicasting by means of a sequencer, one approach is to first forward an operation to the sequencer, which then assigns it a unique number and subsequently multicasts the operation. Mention two alternative approaches, and compare the three solutions.

A: Another approach is to multicast the operation, but defer delivery until the sequencer has subsequently multicast a sequence number for it. The latter happens after the operation has been received by the sequencer. A third approach is to first get a sequence number from the sequencer, and then multicast the operation.

The first approach (send operation to sequencer), involves sending one point-to-point message with the operation, and a multicast message. The second approach requires two multicast messages: one containing the operation, and one containing a sequence number. The third approach, finally, costs one point-to-point message with the sequence number, followed by a multicast message containing the operation.

20. **Q:** A file is replicated on 10 servers. List all the combinations of read quorum and write quorum that are permitted by the voting algorithm.

A: The following possibilities of (read quorum, write quorum) are legal. (1, 10), (2, 9), (3, 8), (4, 7), (5, 6), (6, 5), (7, 4), (8, 3), (9, 2), and (10, 1).

21. **Q:** State-based leases are used to offload a server by letting it allow to keep track of as few clients as needed. Will this approach necessarily lead to better performance?

A: No, for the simple reason that for some clients it would still be better to inform them when updates happened. Not maintaining any state may lead to the situation that these clients will often poll the already busy server.

22. For this exercise, you are to implement a simple system that supports multicast RPC. We assume that there are multiple, replicated servers and that each client communicates with a server through an RPC. However, when dealing with replication, a client will need to send an RPC request to each replica. Program the client such that to the application it appears as if a single RPC is sent. Assume you are replicating for performance, but that servers are susceptible to failures.

SOLUTIONS TO CHAPTER 8 PROBLEMS

1. **Q:** Dependable systems are often required to provide a high degree of security. Why?

A: If, for example, the responses given by servers cannot be trusted to be

correct because some malicious party has tampered with them, it hardly makes sense to talk about a dependable system. Likewise, servers should be able to trust their clients.

2. **Q:** What makes the fail-stop model in the case of crash failures so difficult to implement?

   A: The fact that, in practice, servers simply stop producing output. Detecting that they have actually stopped is difficult. As far as another process can see, the server may just be slow, or communication may (temporarily) be failing.

3. **Q:** Consider a Web browser that returns an outdated cached page instead of a more recent one that had been updated at the server. Is this a failure, and if so, what kind of failure?

   A: Whether or not it is a failure depends on the consistency that was promised to the user. If the browser promises to provide pages that are at most T time units old, it may exhibit performance failures. However, a browser can never live up to such a promise in the Internet. A weaker form of consistency is to provide one of the client-centric models discussed in Chap. 7. In that case, simply returning a page from the cache without checking its consistency may lead to a response failure.

4. **Q:** Can the model of triple modular redundancy described in the text handle Byzantine failures?

   A: Absolutely. The whole discussion assumed that failing elements put out random results, which are the same as Byzantine failures.

5. **Q:** How many failed elements (devices plus voters) can Fig. 8-0 handle? Give an example of the worst case that can be masked.

   A: In each row of circles, at most one element can fail and be masked. Furthermore, one voter in each group can also fail provided it is feeding a faulty element in the next stage. For example, if all six elements at the top of their respective columns all fail, two of the three final outputs will be correct, so we can survive six failures.

6. **Q:** Does TMR generalize to five elements per group instead of three? If so, what properties does it have?

   A: Yes, any odd number can be used. With five elements and five voters, up to two faults per group of devices can be masked.

7. **Q:** For each of the following applications, do you think at-least-once semantics or at most once semantics is best? Discuss.

   (a) Reading and writing files from a file server.
   (b) Compiling a program.
   (c) Remote banking.

A: For (a) and (b), at least once is best. There is no harm trying over and over. For (c), it is best to give it only one try. If that fails, the user will have to intervene to clean up the mess.

8. **Q:** With asynchronous RPCs, a client is blocked until its request has been *accepted* by the server. To what extent do failures affect the semantics of asynchronous RPCs?

   A: The semantics are generally affected in the same way as ordinary RPCs. A difference lies in the fact that the server will not be processing the request while the client is blocked, which introduces problems when the client crashes in the meantime. Instead, the server simply does its work, and attempts to contact the client later on, if necessary.

9. **Q:** Give an example in which group communication requires no message ordering at all.

   A: Multicasting images in small fragments, where each fragment contains the (x, y) coordinate as part of its data. Likewise, sending the pages of a book, with each page being numbered.

10. **Q:** In reliable multicasting, is it always necessary that the communication layer keeps a copy of a message for retransmission purposes?

    A: No. In many cases, such as when transferring files, it is necessary only that the data is still available at the application level. There is no need that the communication layer maintains its own copy.

11. **Q:** To what extent is scalability of atomic multicasting important?

    A: It really depends on how many processes are contained in a group. The important thing to note is, that if processes are replicated for fault tolerance, having only a few replicas may be enough. In that case, scalability is hardly an issue. When groups of different processes are formed, scalability may become an issue. When replicating for performance, atomic multicasting itself may be overdone.

12. **Q:** In the text, we suggest that atomic multicasting can save the day when it comes to performing updates on an agreed set of processes. To what extent can we guarantee that each update is actually performed?

    A: We cannot give such guarantees, similar to guaranteeing that a server has actually performed an operation after having sent an acknowledgement to the client. However, the degree of fault tolerance is improved by using atomic multicasting schemes, and makes developing fault-tolerant systems easier.

13. **Q:** Virtual synchrony is analogous to weak consistency in distributed data stores, with group view changes acting as synchronization points. In this context, what would be the analog of strong consistency?

    A: The synchronization resulting from individual multicasts, be they totally,

causally, or FIFO ordered. Note that view changes take place as special multi-cast messages, which are required to be properly ordered as well.

14. **Q:** What are the permissible delivery orderings for the combination of FIFO and total-ordered multicasting in Fig. 8-0?

    A: There are six orderings possible:

Order 1	Order 2	Order 3	Order 4	Order 5	Order 6
m1	m1	m1	m3	m3	m3
m2	m3	m3	m1	m1	m4
m3	m2	m4	m2	m4	m1
m4	m4	m2	m4	m2	m2

15. **Q:** Adapt the protocol for installing a next view G_{i+1} in the case of virtual synchrony so that it can tolerate process failures.

    A: When a process P receives G_{i+k}, it first forwards a copy of any unstable message it has, regardless to which previous view it belonged, to every process in G_{i+k}, followed by a flush message for G_{i+k}. It can then safely mark the message as being stable.

    If a process Q receives a message m that was sent in G_j (with $j < i + k$), it discards it when Q was never in G_j. If the most recently installed view at Q is G_l with $l > j$, message m is also discarded (it is a duplicate). If $l = j$ and m had not yet been received, process Q delivers m taking any additional message ordering constraints into account. Finally, if $l < j$, message m is simply stored in the communication layer until G_j has been installed.

16. **Q:** In the two-phase commit protocol, why can blocking never be completely eliminated, even when the participants elect a new coordinator?

    A: After the election, the new coordinator may crash as well. In this case, the remaining participants can also not reach a final decision, because this requires the vote from the newly elected coordinator, just as before.

17. **Q:** In our explanation of three-phase commit, it appears that committing a transaction is based on majority voting. Is this true?

    A: Absolutely not. The point to note is that a recovering process that could not take part in the final decision as taken by the other process, will recover to a state that is consistent with the final choice made by the others.

18. **Q:** In a piecewise deterministic execution model, is it sufficient to log only messages, or do we need to log other events as well?

    A: More logging is generally needed: it concerns *all* nondeterministic events, including local I/O and, in particular, system calls.

19. **Q:** Explain how the write-ahead log in distributed transactions can be used to recover from failures.

A: The log contains a record for each read and write operation that took place within the transaction. When a failure occurs, the log can be replayed to the last recorded operation. Replaying the log is effectively the opposite from rolling back, which happens when the transaction needed to be aborted.

20. **Q:** Does a stateless server need to take checkpoints?

    A: It depends on what the server does. For example, a database server that has been handed a complete transaction will maintain a log to be able to redo its operations when recovering. However, there is no need to take checkpoints for the sake of the state of the distributed system. Checkpointing is done only for local recovery.

21. **Q:** Receiver-based message logging is generally considered better than sender-based logging. Why?

    A: The main reason is that recovery is entirely local. In sender-based logging, a recovering process will have to contact its senders to retransmit their messages.

SOLUTIONS TO CHAPTER 9 PROBLEMS

1. **Q:** Which mechanisms could a distributed system provide as security services to application developers that believe only in the end-to-end argument in system's design, as discussed in Chap. 6?

   A: None. Applying the end-to-end argument to security services means that developers will not trust anything that is not provided by their own applications. In effect, the distributed system as a whole is considered to be untrusted.

2. **Q:** In the RISSC approach, can all security be concentrated on secure servers or not?

   A: No, we still need to make sure that the local operating systems and communication between clients and servers are secure.

3. **Q:** Suppose you were asked to develop a distributed application that would allow teachers to set up exams. Give at least three statements that would be part of the security policy for such an application.

   A: Obvious requirements would include that students should not be able to access exams before a specific time. Also, any teacher accessing an exam before the actual examination date should be authenticated. Also, there may be a restricted group of people that should be given read access to any exam in preparation, whereas only the responsible teacher should be given full access.

4. **Q:** Would it be safe to join message 3 and message 4 in the authentication protocol shown in Fig. 9-0, into $K_{A,B}(R_B, R_A)$?

A: Yes, there is no reason why the challenge sent by Alice for Bob cannot be sent in the same message.

5. **Q:** Why is it not necessary in Fig. 9-0 for the KDC to know for sure it was talking to Alice when it receives a request for a secret key that Alice can share with Bob?

   A: Suppose that Chuck had sent the message "I'm Alice and I want to talk to Bob." The KDC would just return $K_{A,KDC}(K_{A,B})$ which can be decrypted only by Alice because she is the only other entity holding the secret key $K_{A,KDC}$.

6. **Q:** What is wrong in implementing a nonce as a timestamp?

   A: Although a timestamp is used only once, it is far from being random. Implementations of security protocols exist that use timestamps as nonces, and which have been succusfully attacked by exploiting the nonrandomness of the nonces.

7. **Q:** In message 2 of the Needham-Schroeder authentication protocol, the ticket is encrypted with the secret key shared between Alice and the KDC. Is this encryption necessary?

   A: No. Because Bob is the only one who can decrypt the ticket, it might as well have been sent as plaintext.

8. **Q:** Can we safely adapt the authentication protocol shown in Fig. 9-0 such that message 3 consists only of R_B?

   A: In principle, if R_B is never used again, then returning it unencrypted should be enough. However, such randomness is seldom found. Therefore, by encrypting R_B, it becomes much more difficult for Chuck to break in and forge message 3.

9. **Q:** Devise a simple authentication protocol using signatures in a public-key cryptosystem.

   A: If Alice wants to authenticate Bob, she sends Bob a challenge R. Bob will be requested to return $K_B^-(R)$, that is, place his signature under R. If Alice is confident that she has Bob's public key, decrypting the response back to R should be enough for her to know she is indeed talking to Bob.

10. **Q:** Assume Alice wants to send a message m to Bob. Instead of encrypting m with Bob's public key K_B^+, she generates a session key $K_{A,B}$ and then sends $[K_{A,B}(m), K_B^+(K_{A,B})]$. Why is this scheme generally better? (*Hint:* consider performance issues.)

    A: The session key has a short, fixed length. In contrast, the message m may be of arbitrary length. Consequently, the combination of using a session key and applying public-key cryptography to a short message will generally provide much better performance than using only a public key on a large message.

11. Q: What is the role of the timestamp in message 6 in Fig. 9-0, and why does it need to be encrypted?

A: The timestamp is used to protect against replays. By encrypting it, it becomes impossible to replay message 6 with a later timestamp. This example illustrates a general application of timestamps in cryptographic protocols.

12. Q: Complete Fig. 9-0 by adding the communication for authentication between Alice and Bob.

A: Alice sends to Bob the message $M = [K_{B,AS}(A, K_{A,B}), K_{A,B}(t)]$, where $K_{B,AS}$ is the secret key shared between Bob and the AS. At that point, Bob knows he is talking to Alice. By responding with $K_{A,B}(t + 1)$, Bob proves to Alice that he is indeed Bob.

13. Q: How can role changes be expressed in an access control matrix?

A: Roles, or protection domains in general, can be viewed as objects with basically a single operation: enter. Whether or not this operation can be called depends on the role from which the request is issued. More sophisticated approaches are also possible, for example, by allowing changes back to previous roles.

14. Q: How are ACLs implemented in a UNIX file system?

A: Each file has three associated entries: one for the owner, one for a group that is associated with the file, and one for everyone else. For each entry, the access rights can essentially be specified as *read*, *write*, *execute*.

15. Q: How can an organization enforce the use of a Web proxy gateway and prevent its users to directly access external Web servers?

A: One way is to use a packet-filtering gateway that discards all outgoing traffic except that directed to a few, well-known hosts. Web traffic is accepted provided it is targeted to the company's Web proxy.

16. Q: Referring to Fig. 9-0, to what extent does the use of Java object references as capabilities actually depend on the Java language?

A: It is independent of the Java language: references to secured objects still need to be handed out during runtime and cannot be simply constructed. Java helps by catching the construction of such references during compile time.

17. Q: Name three problems that will be encountered when developers of interfaces to local resources are required to insert calls to enable and disable privileges to protect against unauthorized access by mobile programs as explained in the text.

A: An important one is that no thread switching may occur when a local resource is called. A thread switch could transfer the enabled privileges to another thread that is not authorized to access the resource. Another problem

occurs when another local resource needs to be called before the current invocation is finished. In effect, the privileges are carried to the second resource, while it may happen that the caller is actually not trusted to access that second resource. A third problem is that explicitly inserting calls to enable and disable privileges is suspect to programming errors, rendering the mechanism useless.

18. **Q:** Name a few advantages and disadvantages of using centralized servers for key management.

    A: An obvious advantage is simplicity. For example, by having N clients share a key with only a centralized server, we need to maintain only N keys. Pairwise sharing of keys would add up to $N(N-1)/2$ keys. Also, using a centralized server allows efficient storage and maintenance facilities at a single site. Potential disadvantages include the server becoming a bottleneck with respect to performance as well as availability. Also, if the server is compromised, new keys will need to be established.

19. **Q:** The Diffie-Hellman key-exchange protocol can also be used to establish a shared secret key between three parties. Explain how.

    A: Suppose Alice, Bob, and Chuck want to set up a shared secret key based on the two publicly known large primes n and g. Alice has her own secret large number x, Bob has y, and Chuck has z. Alice sends $g^x \bmod n$ to Bob; Bob sends $g^y \bmod n$ to Chuck; and Chuck sends $g^z \bmod n$ to Alice. Alice can now compute $g^{xz} \bmod n$, which she sends to Bob. Bob, in turn, can then compute $g^{xyz} \bmod n$. Likewise, after receiving $g^x \bmod n$ from Alice, Bob can compute $g^{xy} \bmod n$, which he sends to Chuck. Chuck can then compute $g^{xyz} \bmod n$. Similarly, after Chuck receives $g^y \bmod n$ from Bob, he computes $g^{yz} \bmod n$ and sends that to Alice so that she can compute $g^{xyz} \bmod n$.

20. **Q:** There is no authentication in the Diffie-Hellman key-exchange protocol. By exploiting this property, a malicious third party, Chuck, can easily break into the key exchange taking place between Alice and Bob, and subsequently ruin the security. Explain how this would work.

    A: Assume Alice and Bob are using the publicly known values n and g. When Alice sends $g^x \bmod n$ to Bob, Chuck need simply intercept that message, return his own message $g^z \bmod n$, and thus make Alice believe she is talking to Bob. After intercepting Alice's message, he sends $g^z \bmod n$ to Bob, from which he can expect $g^y \bmod n$ as reply. Chuck is now in the middle.

21. **Q:** Give a straightforward way how capabilities in Amoeba can be revoked.

    A: The object's owner simply requests that the server discard all registered *(rights, check)*-pairs for that object. A disadvantage is that *all* capabilities are revoked. It is difficult to revoke a capability handed out to a specific process.

22. **Q:** Does it make sense to restrict the lifetime of a session key? If so, give an example how that could be established.

    A: Session keys should always have a restricted lifetime as they are easier to break than other types of cryptographic keys. The way to restrict their lifetime is to send along the expiration time when the key is generated and distributed. This approach is followed, for example, in SESAME.

23. Install and configure a Kerberos v5 environment for a distributed system consisting of three different machines. One of these machines should be running the KDC. Make sure you can setup a (Kerberos) telnet connection between any two machines, but making use of only a single registered password at the KDC. Many of the details on running Kerberos are explained in (Garman, 2003).

SOLUTIONS TO CHAPTER 10 PROBLEMS

1. **Q:** We made a distinction between remote objects and distributed objects. What is the difference?

    A: A remote object is an object that is hosted by a single server, but whose methods can be invoked by remote clients. In contrast, a distributed object is one whose state may be physically distributd across different servers. An example of distributed objects are those provided by Globe. Most object-based systems, however, support only remote objects.

2. **Q:** Why is it useful to define the interfaces of an object in an Interface Definition Language?

    A: There are several reasons. First, from a software-engineering point of view, having precise and unambiguous interface definitions is important for understanding and maintaining objects. Furthermore, IDL-based definitions come in handy for generating stubs. Finally, and related to the latter, if an IDL definition has been parsed and stored, supporting dynamic invocations becomes easier, as the client-side proxy can be automatically constructed at runtime from an interface definition.

3. **Q:** Some implementations of distributed-object middleware systems are entirely based on dynamic method invocations. Even static invocations are compiled to dynamic ones. What is the benefit of this approach?

    A: Realizing that an implementation of dynamic invocations can handle *all* invocations, static ones become just a special case. The advantage is that only a single mechanism needs to be implemented. A possible disadvantage is that performance is not always as optimal as it could be had we analyzed the static invocation.

4. Q: Outline a simple protocol that implements at-most-once semantics for an object invocation.

A: A simple solution is to let a proxy retransmit an invocation request, but telling the server explicitly that it is a retransmission. In that case, the server may decide to ignore the request and return an error to the client. In a more sophisticated solution, the server can cache results of previous invocations and check whether it still has those results when receiving a retransmitted request.

5. Q: Should the client and server-side objects for asynchronous method invocation be persistent?

A: In general, they should be persistent, allowing the client or server to shut down and later restart and fetch the objects from disk. However, in theory, there is no hard reason to demand that these objects should be persistent.

6. Q: In the text, we mentioned that an implementation of CORBA's asynchronous method invocation do not affect the server-side implementation of an object. Explain why this is the case.

A: The important issue is that the client-side runtime system handles all the calls to the server. In particular, the RTS can do a synchronous call to the server, possibly having to wait a long time before an answer is returned. At that moment, it does an upcall to the client application. Likewise, the method invocation can be forwarded to a message router, where eventually, the targeted object server is simply called. Again, it is the communication subsystem that handles the asynchronous nature of the invocation.

7. Q: Give an example in which the (inadvertent) use of callback mechanisms can easily lead to an unwanted situation.

A: If a callback leads to another invocation on the same object, a deadlock may arise if locks are needed to protect shared resources. Situations as these are very hard to control in a general way.

8. Q: Is it possible for an object to have more than one servant?

A: Yes. Recall that an object can be virtually anything. Now consider an object that consists of some data that is stored somewhere in a database, along perhaps with some procedures for manipulating its data. There is no reason why we cannot provide two different access points (by means of servants) to that data. Of course, the object will, in this case, also have two different identifiers.

9. Q: Is it possible to have system-specific implementations of CORBA object references while still being able to exchange references with other CORBA-based systems?

A: Yes, this what the IORs are for. The only issue that is important, is that objects that can be referenced from outside, can be represented by an IOR. In that case, a specific CORBA system will need to provide a gateway for

translating such references into its internal and specific implementation.

10. **Q:** How can we authenticate the contact addresses returned by a lookup service for secure Globe objects?

    A: Simply have each contact address be signed using the object's private key. In that way, we establish a secure binding between an object's identifier, its public key, and its contact addresses.

11. **Q:** What is the key difference between object references in CORBA and those in Globe?

    A: In CORBA, an object reference is essentially what is referred to as a contact address in Globe. In Globe, object references are location independent, meaning that no information whatsoever concerning the current server hosting a (distributed) object is contained in the reference. This information is always part of a CORBA IOR.

12. **Q:** Consider Globe. Outline a simple protocol by which a secure channel is set up between a user proxy (who has access to the Alice's private key) and a replica that we know for certain can execute a given method.

    A: Suppose we found replica R. We can execute the following steps. (1) The user proxy sends the user certificate to the replica. (2) The replica responds with its replica certificate, along with a nonce N_R. (3) The user proxy responds by returning $K^-_{Alice}(N_R, N_{Alice})$. This message will allow R to verify that Alice is on the other end of the channel. It responds with $K^-_R(N_{Alice})$ allowing Alice to verify the identity of R. Note that extra information, like session keys, can be exchanged as well.

13. **Q:** Give an example implementation of an object reference that allows a client to bind to a transient remote object.

    A: Using Java, we can express such an implementation as the following class:

```
public class Object_reference {
    InetAddress  server_address;    // network address of object's server
    int          server_endpoint;   // endpoint to which server is listening
    int          object_identifier; // identifier for this object
    URL          client_code;       // (remote) file containing client-side stub
    byte[]       init_data;         // possible additional initialization data
}
```

The object reference should at least contain the transport-level address of the server where the object resides. We also need an object identifier as the server may contain several objects. In our implementation, we use a URL to refer to a (remote) file containing all the necessary client-side code. A generic array of bytes is used to contain further initialization data for that code. An alternative implementation would have been to directly put the client-code into the

reference instead of a URL. This approach is followed, for example, in Java RMI where proxies are passed as reference.

14. **Q:** Java and other languages support exceptions, which are raised when an error occurs. How would you implement exceptions in RPCs and RMIs?

    A: Because exceptions are initially raised at the server side, the server stub can do nothing else but catch the exception and marshal it as a special error response back to the client. The client stub, on the other hand, will have to unmarshal the message and raise the same exception if it wants to keep access to the server transparent. Consequently, exceptions now also need to be described in an interface definition language.

15. **Q:** How would you incorporate persistent asynchronous communication into a model of communication based on RMIs to remote objects?

    A: An RMI should be asynchronous, that is, no immediate results are expected at invocation time. Moreover, an RMI should be stored at a special server that will forward it to the object as soon as the latter is up and running in an object server.

16. **Q:** Consider a distributed object-based system that supports object replication, in which *all* method invocations are totally ordered. Also, assume that an object invocation is atomic (e.g., because every object is automatically locked when invoked). Does such a system provide entry consistency? What about sequential consistency?

    A: By totally ordering all method invocations, we not only achieve that all method invocations for the same object are carried out in the same order at every one of its replicas, but also that method calls to different objects are carried out in the same order everywhere. As a consequence, such a system provides entry consistency, but in particular sequential consistency.

17. **Q:** Describe a receiver-based scheme for dealing with replicated invocations, as mentioned in the text.

    A: In this case, we let the coordinator of a replicated object imply invoke the replicas of the called object, but let everyone of those replicas return their answer to the coordinator. The latter, in turn, will need to do two things. First, it should be able to detect multiple responses to a previous sent request. Such a detection may require the use of sequence numbers. Second, the coordinator should now about the other replicas in its group, and forward a response to each of them.

SOLUTIONS TO CHAPTER 11 PROBLEMS

1. **Q:** Is a file server implementing NFS version 3 required to be stateless?

   A: No, but the NFS protocols are such that it is possible to provide an implementation using stateless servers.

2. **Q:** Explain whether or not NFS is to be considered a distributed file system.

   A: On can argue that it is not a file system, but merely a protocol that allows local file systems to become accessible to remote clients. In fact, most of the actual file system functionality is not implemented by NFS. Instead, it relies on the Virtual File System interface available in most operating systems.

3. **Q:** Despite that GFS scales well, it could be argued that the master is still a potential bottleneck. What would be a reasonable alternative to replace it?

   A: Considering that master uses a file name to look up a chunk server, we could also implement the master in the form of a DHT-based system and use a hash of the file name as the key to be looked up. In this way, one would obtain a fully decentralized master.

4. **Q:** Using RPC2's side effects is convenient for continuous data streams. Give another example in which it makes sense to use an application-specific protocol next to RPC.

   A: File transfer. Instead of using a pure RPC mechanism, it may be more efficient to transfer very large files using a protocol such as FTP.

5. **Q:** NFS does not provide a global, shared name space. Is there a way to mimic such a name space?

   A: A global name space can easily be mimicked using a local name space for each client that is partially standardized, and letting the automounter mount the necessary directories into that name space.

6. **Q:** Give a simple extension to the NFS lookup operation that would allow iterative name lookup in combination with a server exporting directories that it mounted from another server.

   A: If a lookup operation always returns an identifier for the server from which a directory was mounted, transparent iterative name lookups across multiple servers would be easy. Whenever a server looks up a mount point on which it mounted a remote file system, it simply returns the server's ID for that file system. The client can then automatically mount the directory as well, and contact its associated server to continue name resolution.

7. **Q:** In UNIX-based operating systems, opening a file using a file handle can be done only in the kernel. Give a possible implementation of an NFS file handle for a user-level NFS server for a UNIX system.

   A: The problem to be solved is to return a file handle that will allow the server to open a file using the existing name-based file system interface. One

approach is to encode the file name into the file handle. The obvious drawback is that as soon as the file name changes, its file handles become invalid.

8. **Q:** Using an automounter that installs symbolic links as described in the text makes it harder to hide the fact that mounting is transparent. Why?

   A: After the symbolic link has been followed, the user will not be in the expected directory, but in a subdirectory used by the automounter. In other words, the local home directory for Alice will be */tmp_mnt/home/alice* instead of what she thinks it is, namely */home/alice*. Special support from the operating system or shell is needed to hide this aspect.

9. **Q:** Suppose the current denial state of a file in NFS is *WRITE*. Is it possible that another client can first successfully open that file and then request a write lock?

   A: Yes. If the second client requires read/write access (i.e., value *BOTH*) but no denial (i.e., value *NONE*), it will have been granted access to the file. However, although a write lock may actually be granted, performing an actual write operation will fail. Remember that share reservation is completely independent from locking.

10. **Q:** Taking into account cache coherence as discussed in Chap. 7, which kind of cache-coherence protocol does NFS implement?

    A: Because multiple write operations can take place before cached data is flushed to the server, NFS clients implement a write-back cache.

11. **Q:** Does NFS implement entry consistency?

    A: Yes. Because share reservations and file locking are associated with specific files, and because a client is forced to revalidate a file when opening it and flush it back to the server when closing it, it can be argued that NFS implements entry consistency.

12. **Q:** We stated that NFS implements the remote access model to file handling. It can be argued that it also supports the upload/download model. Explain why.

    A: Because a server can delegate a file to a client, it can effectively support whole-file caching and putting that client in charge of further handling of the file. This model comes close to uploading a file to a client and downloading it later when when the client is finished.

13. **Q:** In NFS, attributes are cached using a write-through cache coherence policy. Is it necessary to forward all attributes changes immediately?

    A: No. For example, when appending data to a file, the server does not really need to be informed immediately. Such information may possibly be passed on when the client flushes its cache to the server.

14. **Q:** What calling semantics does RPC2 provide in the presence of failures?

    A: Considering that the client will be reported an error when an invocation does not complete, RPC2 provides at-most-once semantics.

15. **Q:** Explain how Coda solves read-write conflicts on a file that is shared between multiple readers and only a single writer.

    A: The problem is solved by "defining it away." The semantics of transactions underlying file sharing in Coda, permit treating all readers as accessing the shared file before the writer opened it. Note that read-write conflicts within a specific time interval cannot be solved in this way.

16. **Q:** Using self-certifying path names, is a client always ensured it is communicating with a nonmalicious server?

    A: No. SFS does not solve naming problems. In essence, a client will have to trust that the server named in the path can actually be trusted. In other words, a client will have to put its trust in the name and name resolution process. It may very well be the case that a malicious SFS server is spoofing another server by using its IP address and passing the other server's public key.

17. One of the easiest ways for building a UNIX-based distributed system, is to couple a number of machines by means of NFS. For this assignment, you are to connect two file systems on different computers by means of NFS. In particular, install an NFS server on one machine such that various parts of its file system are automatically mounted when the first machine boots.

18. To integrate UNIX-based machines with Windows clients, one can make use of Samba servers. Extend the previous assignment by making a UNIX-based system available to a Windows client, by installing and configuring a Samba server. At the same time, the file system should remain accessible through NFS.

SOLUTIONS TO CHAPTER 12 PROBLEMS

1. **Q:** To what extent is e-mail part of the Web's document model?

    A: E-mail is not part of the document model underlying the Web, but rather a separate system that has been integrated with hypertext documents by means of client-side software only. For example, most Web browsers provide additional support for handling e-mail, but the actual e-mail messages are not related to hypertext documents in any way.

2. **Q:** In many cases, Web sites are designed to be accessed by users. However, when it comes to Web services, we see that Web sites become dependent on each other. Considering the three-tiered architecture of Fig. 12-0, where would you expect to see the dependency occur?

A: There are two places: from the Web server to an externally available service, or from the application (i.e., CGI) server to an external service. Considering that the Web server is acting as a relatively simple front end, and that most of the complexity for implementing the service offered by a Web site is contained in the application server, we generally see that the application server makes use of other Web services.

3. **Q:** The Web uses a file-based approach to documents by which a client first fetches a file before it is opened and displayed. What is the consequence of this approach for multimedia files?

   A: One of the problems is that in many cases such files are first fetched and stored locally before that can be opened. What is needed, however, is to keep the file at the server and set up a data stream to the client. This approach is supported in the Web, but requires special solutions at both the client and the server.

4. **Q:** One could argue that from an technological point of view Web services do not address any new issues. What is the compelling argument to consider Web services important?

   A: What many people underestimate is the inoperability between systems from different manufacturers in addition to inoperability between systems at different organizations. Web services addresses these differences by essentially prescribing how such systems can be used together. Although this may not be very innovative, it is definitely a challenging problem to come up with the right set of standards that fit the many business processes that are deployed in practice.

5. **Q:** What would be the main advantage of using the distributed server discussed in Chap. 0 to implement a Web server cluster, in comparison to the way the such clusters are organized as shown in Fig. 12-0. What is an obvious disadvantage?

   A: The main advantage is that a client could communicate directly with its assigned Web server, which may show to be beneficial when the cluster is dispersed across different networks (which may happen when dealing with flash crowds). However, the more traditional design relies on well-deployed technology (IPv4) and is also seemingly simpler.

6. **Q:** Why do persistent connections generally improve performance compared to nonpersistent connections?

   A: The real gain is in avoiding connection setup, which requires a 3-way handshake in the case of TCP.

7. **Q:** SOAP is often said to adhere to RPC semantics. Is this really true?

   A: The answer is simply "no." SOAP fundamentally adheres a two-way asynchronous message-passing model, in which a request is explicitly formulated

in terms of a message sent to a server, whereas another message is sent back as a response. SOAP does not prescribe that clients need to wait, not does it provide transparancy that one normally finds with RPC systems.

8. **Q:** Explain the difference between a plug-in, an applet, a servlet, and a CGI program.

A: A plug-in is a piece of code that can be dynamically loaded from a browser's *local* file system. In contrast, an applet is dynamically downloaded from the server to the browser, for which reason security is generally more important. Note that many plug-ins can be dynamically downloaded as well, but generally not without manual interference by the user. A servlet is comparable to a plug-in, but is entirely handled at the server side. A CGI program is a piece of code that runs in a separate process on the same machine as the server.

9. **Q:** In WebDAV, is it sufficient for a client to show only the lock token to the server in order to obtain write permissions?

A: No, the client should also identify itself as the rightful owner of the token. For this reason, WebDAV not only hands over a token to a client, but also registers which client has the token.

10. **Q:** Instead of letting a Web proxy compute an expiration time for a document, a server could do this instead. What would be the benefit of such an approach?

A: An important benefit would be that the expiration time is consistent in a hierarchy of caches. For example, if a browser cache computes a longer expiration time than its associated Web proxy, this would mean that one user would get to see a possibly stale document, while other users that access the same proxy are returned a fresher version.

11. **Q:** With Web pages becoming highly personalized (because they can be dynamically generated on a per-client basis), one could argue that Web caches will soon all be obsolete. Yet, this is most likely not going to happen soon. Explain why.

A: The answer lies in the fact that although pages are highly personalized, the various elements of which they are made can often be effectively shared by many different users. For example, a Web page may consist of various images that are used for many different clients. Such elements are good candidates to keep in a local cache, also because they generally do not change often.

12. **Q:** Does the Akamai CDN follow a pull-based or push-based distribution protocol?

A: Because replicas are fetched on demand after a client has been redirected to a CDN server, Akamai is seen to follow a pull-based protocol.

13. **Q:** Outline a simple scheme by which an Akamai CDN server can find out that a cached embedded document is stale without checking the document's validity at the original server.

    A: A simple approach followed by Akamai is to include a hash value of the embedded document in its modified URL. It is important to realize that any client will always retrieve all modified URLs when fetching the main document. Subsequent lookups will eventually lead to a CDN server, which can then conclude that a cached document is no longer valid by comparing hash values in the modified URL of the requested and the cached document, respectively.

14. **Q:** Would it make sense to associate a replication strategy with each Web document separately, as opposed to using one or only a few global strategies?

    A: Probably, considering the wide variety of usage patterns for Web documents. Many documents such as personal home pages, are hardly ever updated, and if so, they are updated by only a single person, making these documents suitable for caching and replication. Dynamically generated documents containing timely information require a different strategy, especially if they need to be replicated for performance. Note that the current Web can hardly differentiate such documents.

15. **Q:** Assume that a nonreplicated document of size s bytes is requested r times per second. If the document is replicated to k different servers, and assuming that updates are propagated separately to each replica, when will replication be cheaper than when the document is not replicated?

    A: The consumed bandwidth in the nonreplicated case is $r \times s$, assuming that the requests are independent. Suppose now that document is updated w times per second and replicated to k servers. Updates are propagated independently, so that we have a total consumed bandwidth for propagating an update as $k \times w \times s$ bytes. Replication will be cheaper only if $k \times w \times s < r \times s$.

16. **Q:** Consider a Web site experiencing a flash crowd. What could be an appropriate measure to take in order to ensure that clients are still serviced well?

    A: One possible solution is to have the Web site replicated and redirect clients to the replicas. This will still put a burden on the main site as it needs to perform the redirections, but in practice this approach works reasonably well. The problem is that the Web site will need to be replicated in advance. As an alternative, if flash crowds are predictable, it may be possible to perform replication on demand, provided the resources to do so can be instantly claimed. The latter is generally the case in CDNs.

17. **Q:** There are, in principle, three different techniques for redirecting clients to servers: TCP handoff, DNS-based redirection, and HTTP-based redirection. What are the main advantages and disadvantages of each technique?

A: Roughly speaking, we have that TCP handoff is highly transparent to clients: they need not notice that there connection is being handled by another server. However, TCP handoff clearly has scalability problems: it will generally work only for Web server clusters. DNS-based redirection is also transparent, but works only for an entire site. Redirection on a per-page basis is not possible. In contrast, HTTP-based redirection can operate at this finer granularity, but has the disadvantage that a client will notice to which page it is being redirected. The latter may not be what it wanted when clients store these redirects such as in the case of maintaining bookmarks in Web browsers.

18. **Q:** Give an example in which a query containment check as performed by an edge server supporting content-aware caching will return successfully.

    A: Consider the two SQL queries Q1 "select * from booklist with price < 10" and Q2 "select * from booklist with price < 20." In this example, the answer for Q1 is contained in that of Q2. As a consequence, if the server had cached the answer for Q2, it should be able to successfully perform Q1 locally.

19. Set up a simple Web-based system by installing and configuring the Apache Web server for your local machine such that it can be accessed from a local browser. If you have multiple computers in a local-area network, make sure that the server can be accessed from any browser on that network.

20. WebDAV is supported by the Apache Web server and allows multiple users to share files for reading and writing over the Internet. Install and configure Apache for a WebDAV-enabled directory in a local-area network. Test the configuration by using a WebDAV client.

SOLUTIONS TO CHAPTER 13 PROBLEMS

1. **Q:** What type of coordination model would you classify the message-queuing systems discussed in Chap. 0?

   A: Considering that in message-queuing systems processes are temporally uncoupled, but will otherwise have to agree on message format and destination queues, they classify as mailbox coordination models.

2. **Q:** Outline an implementation of a publish/subscribe system based on a message-queuing system like that of IBM WebSphere.

   A: Such an implementation can be accomplished by using a message broker. All publish/subscribe messages are published to the broker. Subscribers pass their subscriptions to the broker as well, which will then take care to forward published messages to the appropriate subscribers. Note that the broker makes use of the underlying message-queuing network.

3. Q: Explain why decentralized coordination-based systems have inherent scalability problems.

A: The crux of the matter lies in the fact that we need to do decentralized matching of published data and subscriptions. Sometimes, building an index is possible, such as with subject-based systems. In those cases, we can use DHT-based systems to distribute the index among the nodes in the peer-to-peer system. If building such an index is not possible, as in the case of content-based matching, we essentially need to do a global search to find subscriptions that match published data. Such a search does not scale well.

4. Q: To what is a subject name in TIB/Rendezvous actually resolved, and how does name resolution take place?

A: A name is resolved to the current group of subscribers. Name resolution takes place by properly routing message to those subscribers. In TIB/Rendezvous, routing takes place through a combination of multicasting and filtering messages at rendezvous and router daemons.

5. Q: Outline a simple implementation for totally-ordered message delivery in a TIB/Rendezvous system.

A: Use a sequencer to which all messages are sent. Subscribers pass subscriptions to this sequencer. The FIFO-ordered message delivery of the TIB/Rendezvous system will then guarantee that all messages multicast by the sequencer are delivered to every subscriber in the same order.

6. Q: In content-based routing such as used in the Siena system, which we described in the text, we may be confronted with a serious management problem. Which problem is that?

A: The issue here, is that these type of systems assume that an overlay of application-level routers has already been established. In practice, this means that an organization will need to manually configure the overlay network and ensure that it is kept alive.

7. Q: Assume a process is replicated in a TIB/Rendezvous system. Give two solutions to avoid so that messages from this replicated process are not published more than once.

A: A first solution is to attach a message identifier to each published message, and to let subscribers discard duplicates by taking a look at the identifiers. The main drawback of this approach is the waste of network bandwidth. Another solution is to assign a coordinator among the replicas, and let only the coordinator actually publish messages. This solution is similar to assigning a coordinator in the case of invocations with replicated distributed objects.

8. Q: To what extent do we need totally-ordered multicasting when processes are replicated in a TIB/Rendezvous system?

A: Assuming that duplicate messages are either detected by subscribers, or avoided altogether, total ordering is not an issue at all. In this case, the FIFO-ordering delivery semantics are sufficient to let subscribers process the messages in the order they were published by the replicated process.

9. **Q:** Describe a simple scheme for PGM that will allow receivers to detect missing messages, even the last one in a series.

   A: A simple scheme is to use sequence numbers. Whenever a sender has no more data to send, it should announce by multicasting a special control message. If that control message is lost, a receiver will start complaining in the usual way. The sender can then simply retransmit the control message. In essence, this is also the solution adopted in PGM.

10. **Q:** How could a coordination model based on generative communication be implemented in TIB/Rendezvous?

    A: This is actually not very difficult. What makes TIB/Rendezvous different from, for example, the JavaSpaces in Jini, is that processes are still temporally coupled. If published messages are temporarily stored, it would be possible for subscribers to read them even when the publisher of messages no longer exists. What is needed is for each subscriber to record a published message that it has already read, so that receiving duplicates can be avoided.

11. **Q:** A lease period in Jini is always specified as a duration and not as an absolute time at which the lease expires. Why is this done?

    A: Especially in wide-area systems, it may happen that clocks on different machines give very different times. In that case, specifying the expiration of a lease as an absolute time is simply too inaccurate as the holder of the lease may have a very different idea when the lease expires than the processes that handed out the lease. With durations, this difference becomes less an issue, provided some guarantees can be given that the transmission time of a lease is relatively low.

12. **Q:** What are the most important scalability problems in Jini?

    A: One important problem is related to the fact the Jini uses a multicast protocol to locate lookup services. In a wide-area system, this protocol will have to be replaced by something different if Jini is to scale to large numbers of users and processes. A second problem is related to matching templates to tuples in JavaSpaces. Again, special measures will be needed to search a JavaSpace that may be potentially distributed across a large-scale network. No efficient solutions to these problems are yet known.

13. **Q:** Consider a distributed implementation of a JavaSpace in which tuples are replicated across several machines. Give a protocol to delete a tuple such that race conditions are avoided when two processes try to delete the same tuple.

    A: Use a two-phase commit protocol. In phase one, the process doing the

delete sends a message to all the JavaSpace servers holding the tuple asking
them to lock the tuple. When all of them are locked, the delete is sent. If a sec-
ond delete happens simultaneously, it can happen that some servers have
locked one tuple and some have locked the other. If a server cannot grant a
request because the tuple is already locked, it sends back a negative acknowl-
edgement, which causes the initiator to abort the delete, unlock all the tuples it
has locked, wait a random time, and try again later.

14. **Q:** Suppose a transaction T in Jini requires a lock on an object that is currently
locked by another transaction T'. Explain what happens.

    A: Transaction T will continue to block until the lock is either released or until
the lease on the transaction expires. In the latter case, transaction T is simply
aborted.

15. **Q:** Suppose a Jini client caches the tuple it obtained from a JavaSpace so that
it can avoid having to go to the JavaSpace the next time. Does this caching
make any sense?

    A: Caching is senseless because the tuple will have been removed from the
JavaSpace when it was returned; it is ready for the client to keep. The main
idea behind caching is to keep data local to avoid another server access. In the
case of Jini, a JavaSpace is often used to explicitly synchronize processes.
Caching does not play a role when process synchronization is explicitly
needed.

16. **Q:** Answer the previous question, but now for the case that a client caches the
results returned by a lookup service.

    A: This is a completely different situation. The lookup service stores informa-
tion on the whereabouts of services. In this case, it may indeed make sense for
a client to cache previously returned results and try to contact the returned ser-
vices before going to the lookup service again.

17. **Q:** Outline a simple implementation of a fault-tolerant JavaSpace.

    A: The simplest approach is to implement a JavaSpace on a single server with
stable storage. Write operations succeed only if the tuple has been safely writ-
ten to storage. In a more advanced setting, a distributed JavaSpace can be used
in which a server group is used to mask process failures. In that case, the
servers may need to follow a two-phase commit protocol for each write opera-
tion.

18. **Q:** In some subject-based publish/subscribe systems, secure solutions are
sought in end-to-end encryption beteen publishers and subscribers. However,
this approach may violate the initial design goals of coordination-based sys-
tems. How?

    A: The problem is that end-to-end encryption requires that a publisher and
subscriber know each other. This is against the initial design goal that

processes should be referentially decoupled.